

Order Database Implementation on Linux using Oracle 19c

Hanif Lumsden

University of Maryland Global Campus

DBST 670 9040 Database Administration

Dr. Igwe

11/15/2022

Table of Contents

Table of Contents	2
Introduction	3
Database Planning	3
Project Plan	3
Scope	5
In-Scope	5
Out-of-Scope	6
Hardware and Software Tools	6
Requirements Definition	6
Business Rules	6
Data Tables	7
Entity Relationship Diagram	7
Capacity Planning	8
Creating the Instance	14
User, Tablespace and Data Definition Language Creation	15
Foreign Keys	19
Creating Partitioned Table	21
Export and Import	23
Loading .ctl and .dat files using SQL*Loader	24
Table Analyze	29
Table Sizes	33
Monitoring the Database Statistics	34
Statistic Interpretation	39
Performance Tuning with Automatic Database Diagnostic Monitor (ADDM)	43
Performance Testing Matrix	56
Conclusion	57
References	57
Index	59
Database Configuration Assistant Wizard Steps	59

Introduction

Order management structures aid businesses in processing transactions made by consumers from various sales channels. In-store exchanges and e-commerce means such as online shopping comprise the transactions of the order database (Amula, 2020). Order database management systems track people, processes, and suppliers entailing the administration of business processes related to the company's order (Amula, 2020, para. 1). The oracle database administrator will locate for their client six preexisting order management tables, and create four relevant ones that will complete the order database. The database will undergo data definition and manipulation language,

Database Planning

Project Plan

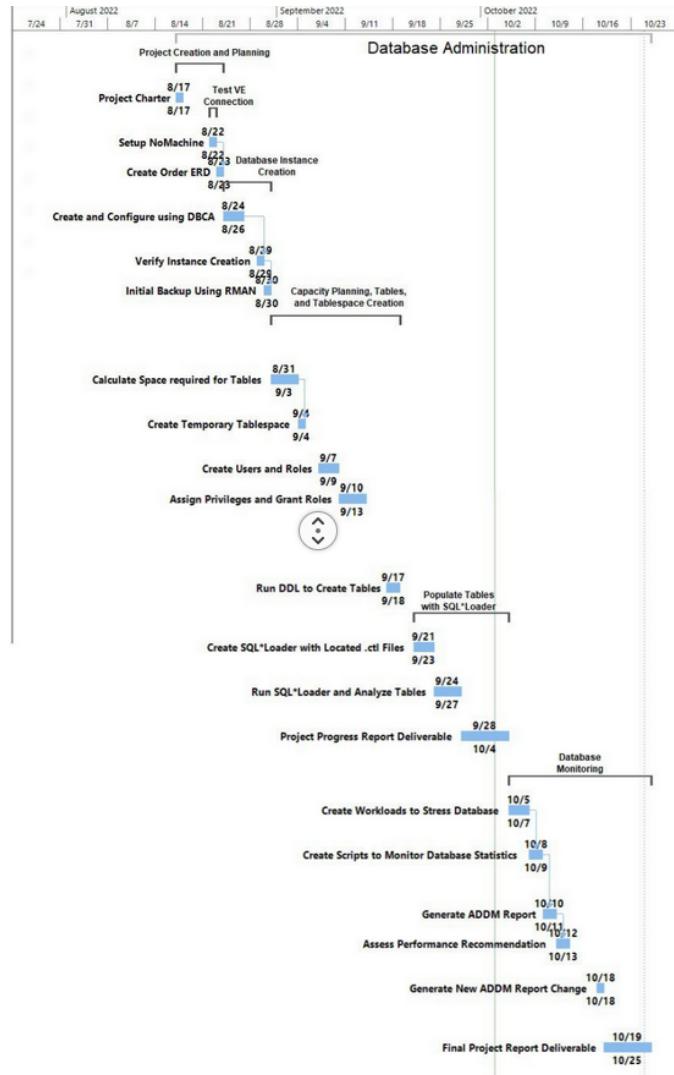
The database administration project plan follows five phases: project creation and planning; database instance creation; capacity planning, tables, and tablespace creation; populate tables with SQL*Loader; and database monitoring. Below is the gantt chart that coincides with the project plan schedule. Copies of the project plan schedule can be accessed [here](#) and gantt chart accessed [here](#).

Figure 2

Project Chart

#	Task Name	Duration	Start	Finish
1	Database Administration	7 days	Wed 8/27/22 8:00 AM	Tue 9/2/22 5:00 PM
2	Project Creation	7 days	Wed 8/27/22 8:00 AM	Tue 9/2/22 5:00 PM
3	Project Charter	1 day	Mon 8/29/22 8:00 AM	Wed 8/30/22 5:00 PM
4	Requirements Gathering	1 day	Mon 8/29/22 8:00 AM	Mon 8/29/22 5:00 PM
5	Database Connection	1 day	Mon 8/29/22 8:00 AM	Mon 8/29/22 5:00 PM
6	Network Configuration	1 day	Tue 8/30/22 8:00 AM	Tue 8/30/22 5:00 PM
7	Database Instance	7 days	Wed 8/30/22 8:00 AM	Tue 9/5/22 5:00 PM
8	Create and Configure Database	3 days	Wed 8/30/22 8:00 AM	Fri 8/31/22 5:00 PM
9	DBCA	1 day	Mon 8/29/22 8:00 AM	Mon 8/29/22 5:00 PM
10	Virtual Database	1 day	Tue 8/30/22 8:00 AM	Tue 8/30/22 5:00 PM
11	Creating Tables	1 day	Tue 8/30/22 8:00 AM	Tue 8/30/22 5:00 PM
12	Creating Tables and Indexes	19 days	Wed 8/31/22 8:00 AM	Sun 9/18/22 5:00 PM
13	Create Schema required for Virtual Database	4 days	Wed 8/31/22 8:00 AM	Sat 9/3/22 5:00 PM
14	Create Temporary Tables	1 day	Thu 8/31/22 8:00 AM	Thu 8/31/22 5:00 PM
15	Create User and Roles	1 days	Wed 8/31/22 8:00 AM	Fri 8/31/22 5:00 PM
16	Grant Privileges and Grant Roles	4 days	Sat 9/3/22 8:00 AM	Tue 9/6/22 5:00 PM
17	Grant Privileges and Create Oracle and Permanent Tables	3 days	Wed 8/31/22 8:00 AM	Fri 8/31/22 5:00 PM
18	Run DDL to Create Oracle and Permanent Tables	2 days	Sat 8/31/22 8:00 AM	Sun 9/3/22 5:00 PM
19	Populate Tables	14 days	Wed 8/31/22 8:00 AM	Tue 9/13/22 5:00 PM
20	Create Tables with Data Located off File	3 days	Wed 8/31/22 8:00 AM	Fri 8/31/22 5:00 PM
21	Run SQL Scripts and Analyze Data	4 days	Sat 8/31/22 8:00 AM	Tue 9/5/22 5:00 PM
22	Project Progress Report	7 days	Wed 8/28/22 8:00 AM	Tue 9/12/22 5:00 PM
23	Deliverable Preparation	21 days	Wed 8/28/22 8:00 AM	Tue 9/27/22 5:00 PM
24	Meeting Charter	2 days	Wed 8/28/22 8:00 AM	Fri 8/30/22 5:00 PM
25	Workload to Identify Delays	3 days	Wed 8/29/22 8:00 AM	Fri 8/31/22 5:00 PM
26	Create Scripts for Database Migration	2 days	Sat 8/30/22 8:00 AM	Sun 8/31/22 5:00 PM
27	Generate ADDM Report	2 days	Mon 8/30/22 8:00 AM	Tue 8/31/22 5:00 PM
28	Addm Analysis	2 days	Wed 8/31/22 8:00 AM	Thu 8/31/22 5:00 PM
29	Recommendation	1 day	Tue 8/30/22 8:00 AM	Tue 8/30/22 5:00 PM
30	Generate ADDM Report Change	1 day	Tue 8/30/22 8:00 AM	Tue 8/30/22 5:00 PM
31	Final Project Report	7 days	Wed 8/29/22 8:00 AM	Tue 9/5/22 5:00 PM
32	Deliverable			

Figure 3*Project Gantt Chart*



Scope

In-Scope

The following are within the scope of the project:

- Assessment of the linux virtual environment.
- Creation of logical and physical models.
- Instance creation and backup of the database using RMAN.
- Data definition and manipulation languages generated and said queries ran into the database.

- Creation of users, roles, indices and tablespaces
- Export and importing data.
- Using SQL*Loader to load the sample order data into the database.
- Verifying objects and analyzing tables using the COMPUTE STATISTICS clause.
- Monitoring database statistics.
- Performance tuning with the Automatic Database Diagnostic Monitor.

Out-of-Scope

The following are out of the scope:

- Populating the database system with actual data.
- Enacting role-based security measures and data redaction.

Hardware and Software Tools

Data Modeler:

- Oracle SQL*Loader

Database Management System

- Oracle Database 19c Enterprise Edition Release 19.0.0.0.0 - Production Version
Oracle SQL Developer 19.3.0.0.0 64-bit running on Windows 10

Hardware and Software

- Intel(R) Core(TM) i5-8400 CPU @ 2.80GHz - 2.81 GHz, 8.00 GB (7.88 GB usable), 64-bit operating system, x64-based processor, Windows 11 Home v. 21H2 OS build 22000.1098

Requirements Definition

Business Rules

- a. A ORDERS can has one ORDER_ITEM.

- b. A ORDER_ITEM contains one to many ORDERS.
- c. A PRODUCTS provides one to many for the SUPPLIER.
- d. A CUSTOMER places one to many ORDERS.
- e. A ORDERS requires one to many PAYMENTS.
- f. A EMPLOYEE receives one to many ORDERS.
- g. A EMPLOYEE works for one or many STORES.
- h. A CUSTOMER belongs to one LOCATION.
- i. A STORE exists in one LOCATION.
- j. A SUPPLIER operates in one LOCATION.
- k. A MANAGER manages one STORE.
- l. A EMPLOYEE reports to one to many MANAGER.
- m. A MANAGER resides in one LOCATION.

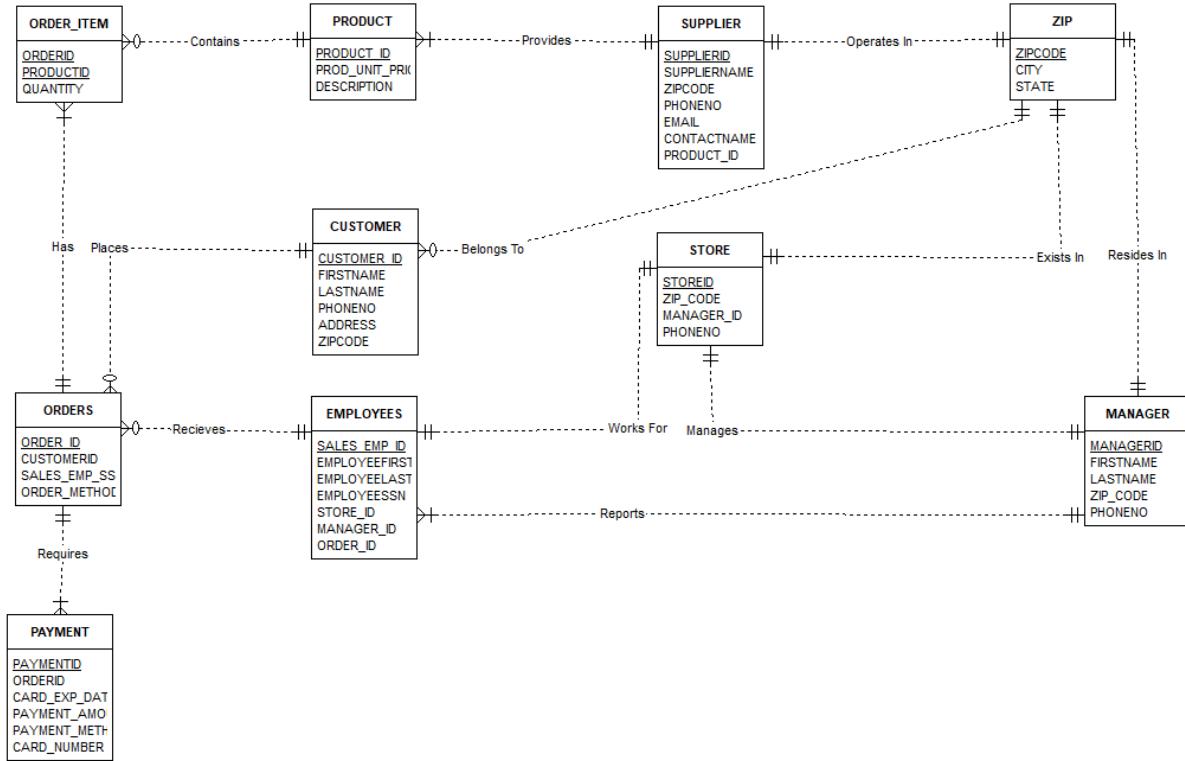
Data Tables

The database will have 10 tables: Orders, Payment, Order_Item, Products, Supplier, Location, Store, Manager, Employees,

Entity Relationship Diagram

Figure 1

Physical Model



Capacity Planning

This statement is referred for capacity calculations. overhead = fixed header + variable transaction header + table directory + row directory.

Knowing that:

Block size is 8192 bytes

overhead = 8192 bytes

fixed header = 57 bytes

transaction header = 23 * initrans value for the table

table directory = 4 (for non-cluster databases)

row directory = 2*R

R = number of rows in the block

Therefore:

$$\text{overhead} = 57 \text{ bytes} + 23*i + 4 \text{ bytes} + 2R$$

$$\text{overhead} = 61 \text{ bytes} + 23*(1) + 2R$$

$$8192 \text{ bytes} = 84 \text{ bytes} + 2R$$

$$8108 \text{ bytes} = 2R$$

$$\text{data space} = \text{available space} - (\text{available space} * \text{pctfree})$$

$$\text{data space} = (8108 \text{ bytes} - 2R) - ((8108-2R)*1/10)$$

$$\text{data space} = 8108 - 2R - 810.8 + .2R$$

$$\text{data space} = 7297.2 - 1.8R$$

We can calculate. Below are assumptions in terms of the columns and their sizes:

Products

$$\text{ProductID} \rightarrow 10 \text{ bytes}$$

$$\text{Prod_Unit_Price} \rightarrow \log(100) / \log(2) / 8 = 1 \text{ byte}$$

$$\text{Description} \rightarrow 100 \text{ bytes}$$

$$\text{Supplier_ID} \rightarrow 10 \text{ bytes}$$

$$(10 + 1 + 100 + 10) \text{ bytes} = 122 \text{ bytes}$$

$$R = (7297.2 - 1.8*R)/122$$

$$1.015R = 59.81$$

$$R = 58.93$$

Order_item

$$\text{order_id} \rightarrow 8 \text{ byte}$$

$$\text{productID} \rightarrow 10 \text{ byte}$$

$$\text{quantity} \rightarrow \log(1000) / \log(2) / 8 = 1 \text{ byte}$$

$(8 + 10 + 1)$ bytes = 19 bytes

$$R = (7297.2 - 1.8 \cdot R) / 19$$

$$1.09R = 384.06$$

$$R = 352.35$$

Payment

PaymentID -> 10 bytes

OrderID -> 10 byte

CardExpDate -> Date = 4 bytes

PaymentAmount -> $\log(100) / \log(2) / 8 = 1$

PaymentMethod -> 10 bytes

$(10 + 10 + 4 + 1 + 10)$ bytes = 35 bytes

$$R = (7297.2 - 1.8 \cdot R) / 35$$

$$1.051R = 208.49$$

$$R = 198.37$$

Orders

OrderID -> 8 byte

CustomerID -> 10 bytes

SalesEmpID -> 8 byte

OrderMethod -> 10 bytes

$(8 + 10 + 8 + 10)$ bytes = 36 bytes

$$R = (7297.2 - 1.8 \cdot R) / 36$$

$$1.05R = 202.7$$

$$R = 193.05$$

Customer

CustomerID -> 10 bytes

FirstName -> 13 bytes

LastName -> 13 bytes

PhoneNo -> 5 byte

LocationID -> 100 bytes

$(10 + 13 + 13 + 5 + 100) \text{ bytes} = 141 \text{ bytes}$

$$R = (7297.2 - 1.8*R)/141$$

$$1.02R = 86.87$$

$$R = 85.17$$

Zip

LocationID -> 10 bytes

ZipCode -> 1 byte

C -> 25 bytes

State -> 25 bytes

Address -> 50 bytes

$(10 + 1 + 25 + 25 + 50) \text{ bytes} = 110 \text{ bytes}$

$$R = (7297.2 - 1.8*R)/110$$

$$1.016R = 66.34$$

$$R = 65.3$$

Supplier

SupplierID -> 10 bytes

SupplierName -> 50 bytes

LocationID -> 10 bytes

PhoneNumber -> 5 byte

Email -> 50 bytes

ContactName - > 50 bytes

$(10 + 50 + 10 + 5 + 50 + 50) \text{ bytes} = 175 \text{ bytes}$

$$R = (7297.2 - 1.8*R)/175$$

$$1.01R = 41.70$$

$$R = 41.3$$

Store

StoreNo -> 1 byte

Address -> 40 bytes

ZipCode -> 1 byte

PhoneNo -> 1 byte

Email -> 30 bytes

ContactName -> 40 bytes

$(1 + 40 + 1 + 1 + 30 + 40) \text{ bytes} = 113 \text{ bytes}$

$$R = (7297.2 - 1.8*R)/113$$

$$1.0159R = 64.57$$

$$R = 63.56$$

Employees

SalesEmpID -> 8 bytes

EmployeeFirstName -> 25 bytes

EmployeeLastName -> 25 bytes

EmployeeSSN -> 5 byte

StoreID -> 10 bytes

ManagerID -> 10 bytes

$(8 + 25 + 25 + 5 + 10 + 10)$ bytes = 83 bytes

$$R = (7297.2 - 1.8 \cdot R) / 83$$

$$1.0217R = 87.92$$

$$R = 86.05$$

Manager

ManagerID -> 10 byte

FirstName -> 25 bytes

LastName -> 25 bytes

LocationID -> 10 bytes

PhoneNumber -> 5 bytes

ContactName -> 40 bytes

$(10 + 25 + 25 + 10 + 5 + 40)$ bytes = 115 bytes

$$R = (7297.2 - 1.8 \cdot R) / 115$$

$$1.0156R = 63.45$$

$$R = 62.47$$

Fetching the .dat files, we are able to see the number of row entries.

```
[oracle@NIXORA01 dir_6tables]$ cat customer.dat|wc -l
1006
[oracle@NIXORA01 dir_6tables]$ cat order_item.dat|wc -l
1102
[oracle@NIXORA01 dir_6tables]$ cat order.dat|wc -l
1103
```

```
[oracle@NIXORA01 dir_6tables]$ cat payment.dat |wc -l
1003
[oracle@NIXORA01 dir_6tables]$ cat product.dat |wc -l
14279
[oracle@NIXORA01 dir_6tables]$ cat hb_zip.dat| wc -l
50002
```

Block calculations are as follows. The manager, employee, store, and supplier tables are user constructed:

ProductBlock = ceiling(14279/58.93) = ceiling(242.30) = 243 blocks = 1990656 bytes

Order_itemBlock = ceiling(1102/352.35) = ceiling(3.127) = 4 blocks = 32768 bytes

PaymentBlock = ceiling(1003/198.37) = ceiling(5.06) = 6 blocks = 49152 bytes

OrdersBlock = ceiling(1103/193.05) = ceiling(5.71) = 6 blocks = 49152 bytes

CustomerBlock = ceiling(1006/85.17) = ceiling(11.81) = 12 blocks = 98304 bytes

ZipBlock = ceiling(50002/65.3) = ceiling(765.72) = 766 blocks = 6275072 bytes

SupplierBlock = ceiling(1000/41.3) = ceiling(24.21) = 24 blocks = 196608 bytes

StoreBlock = ceiling(1000/63.56) = ceiling(15.73) = 16 blocks = 131072 bytes

EmployeesBlock = ceiling(1000/86.05) = ceiling(11.62) = 12 blocks = 98304 bytes

ManagerBlock = ceiling(1000/62.47) = ceiling(16.01) = 16 blocks = 131072 bytes

With the block calculations done, it is compared with the actual table sizes.

Creating the Instance

Logging into the Linux os virtual machine, the user establishes a connection to 10.11.55.52 using the NoMachine software. The following is run in the terminal to add the non-network location connections to the access control list, then the database configuration assistant is started:

```
[StudentFirst@NIXORA01 ~]$ xhost +local:oracle
[StudentFirst@NIXORA01 ~]$ sudo -u oracle $ORACLE_HOME/bin/dbca
```

The user creates the database using the configuration assistant wizard and can be referenced in the index. Initial databases are deleted, and the user goes with advanced configuration. It is an transaction processing Oracle Single Instance Database with the global database name of DBST670. File system location is:

{ORACLE_BASE}/oradata/{DB_UNIQUE_NAME}. Fast recovery storage is 8256 megabytes in the /home area. Listener port 1521 in the Oracle home directory. SGA size is 3000 and PGA size is 1000 using the automatic shared memory management configuration. Processes sizing of 640. Character set is the universal unicode AL32UTF8. Control files are confirmed. Dedicated server mode initialized. RMAN is also initialized to backup the data.

```
[StudentFirst@NIXORA01 ~]$ rman target /
RMAN> backup database plus archivelog;
RMAN> backup archivelog all;
```

User, Tablespace and Data Definition Language Creation

The database is created for the user. Default tablespace is set to USERS and the defeault temporary tables to TEMP. CONNECT and RESOURCE roles are granted as well as the DBA.. An unprivileges user is granted the owner of the tables for the database with the user name ‘hanifdba’ password ‘oracle’..

```
[StudentFirst@NIXORA01 ~]$ sudo -s -u oracle
[oracle@NIXORA01 StudentFirst] . oraenv
ORACLE_SID = [NIXORA01] ? DBST670
The Oracle base remains unchanged with value /u01/app/oracle
[oracle@NIXORA01 StudentFirst]$ sqlplus /nolog

SQL*Plus: Release 19.0.0.9.0 - Production on Thu Sep 15 11:15:03 2022
```

Version 19.3.0.0

Copyright © 1982, 2019 Oracle. All rights reserved.

```
SQL> startup
SQL> connect / as sysdba;
Connected.
SQL> create user hlumsden identified by oracle default tablespace
users temporary tablespace temp;
User created.
SQL> create user hanifdba identified by oracle default tablespace
users temporary tablespace temp quota unlimited on users;
SQL> grant connect, resource to hlumsden;
User created.
SQL> grant hanifdba to dbauser;
Grant succeeded.
SQL> CREATE TABLESPACE "USER_DATA1" LOGGING DATAFILE
'./u01/app/oracle/oradata/DBST670/user_data101.dbf' SIZE 4286K REUSE
AUTOEXTEND ON NEXT 512K MAXSIZE 65536K EXTENT MANAGEMENT LOCAL;
Tablespace created.
SQL> CREATE TABLESPACE "USER_DATA2" LOGGING DATAFILE
'./u01/app/oracle/oradata/DBST670/user_data201.dbf' SIZE 4134K REUSE
AUTOEXTEND ON NEXT 512K MAXSIZE 65536K EXTENT MANAGEMENT LOCAL;
Tablespace created.
SQL> CREATE TABLESPACE "USER_INDX" LOGGING DATAFILE
'./u01/app/oracle/oradata/DBST670/user_idx01.dbf' SIZE 4210K REUSE
AUTOEXTEND ON NEXT 512K MAXSIZE 65536K EXTENT MANAGEMENT LOCAL;
Tablespace created.
SQL> CREATE TEMPORARY TABLESPACE "USER_TEMP" TEMPFILE
'./u01/app/oracle/oradata/DBST670/user_temp01.dbf' SIZE 16384K REUSE
AUTOEXTEND ON NEXT 512K MAXSIZE 65536K EXTENT MANAGEMENT LOCAL;
Tablespace created.
SQL> alter user hanif default tablespace user_data1 temporary
tablespace user_temp;
User altered.
SQL> SHOW USER
USER is "HANIF"
SQL> DROP TABLE customer CASCADE CONSTRAINTS PURGE;
```

```
Table dropped.  
SQL> CREATE TABLE customer (customer_id INTEGER NOT NULL, firstname  
VARCHAR2(15),  
lastname VARCHAR2(15), phoneno CHAR(15), zip_code CHAR(5), address  
CHAR(35)) TABLESPACE  
USER_DATA1;  
Table created.  
SQL> ALTER TABLE customer ADD CONSTRAINT customer_pk PRIMARY KEY  
(customer_id) USING  
INDEX TABLESPACE USER_INDX  
Table altered.  
SQL>  
SQL> DROP TABLE employees CASCADE CONSTRAINTS PURGE;  
Table dropped.  
SQL> CREATE TABLE employees (sales_emp_id INTEGER NOT NULL, store_id  
INTEGER,  
order_id INTEGER, manager_id INTEGER, employeessn VARCHAR2(9),  
employeename  
VARCHAR2(50), employeefirstname VARCHAR2(50)) TABLESPACE USER_DATA1;  
Table created.  
SQL> ALTER TABLE employees ADD CONSTRAINT employees_pk PRIMARY KEY (  
sales_emp_id )  
USING INDEX TABLESPACE USER_INDX;  
Table altered.  
SQL> DROP TABLE manager CASCADE CONSTRAINTS PURGE;  
Table dropped.  
SQL> CREATE TABLE manager (managerid INTEGER NOT NULL, phoneno  
INTEGER, firstname  
VARCHAR2(150), lastname VARCHAR2(150), zip_code CHAR(5)) TABLESPACE  
USER_DATA1;  
Table created.  
SQL> ALTER TABLE manager ADD CONSTRAINT manager_pk PRIMARY KEY (  
managerid ) USING  
INDEX TABLESPACE USER_INDX;  
Table altered.  
SQL> DROP TABLE order_item CASCADE CONSTRAINTS PURGE;  
Table dropped.  
SQL> CREATE TABLE order_item (orderid INTEGER NOT NULL, productid  
INTEGER NOT NULL,
```

```
quantity INTEGER) TABLESPACE USER_DATA1;
Table created.
SQL> ALTER TABLE order_item ADD CONSTRAINT order_item_pk PRIMARY KEY
(
orderid,productid ) USING INDEX TABLESPACE USER_INDX;
Table altered.
SQL> DROP TABLE orders CASCADE CONSTRAINTS PURGE;
Table dropped
SQL> CREATE TABLE orders (order_id INTEGER NOT NULL, customerid
INTEGER, order_method
CHAR(6), sales_emp_ssn INTEGER) TABLESPACE USER_DATA1;
Table created.
SQL> ALTER TABLE orders ADD CONSTRAINT orders_pk PRIMARY KEY (
order_id ) USING INDEX
TABLESPACE USER_INDX;
Table altered.
SQL> DROP TABLE PAYMENTS CASCADE CONSTRAINTS PURGE;
Table dropped.
SQL> CREATE TABLE payment ( paymentid INTEGER NOT NULL, orderid
INTEGER,
payment_amount INTEGER, payment_method VARCHAR2(5), card_number
INTEGER, card_exp_date
DATE) TABLESPACE USER_DATA2;
Table created.
SQL> ALTER TABLE payment ADD CONSTRAINT payment_pk PRIMARY KEY (
paymentid ) USING
INDEX TABLESPACE USER_INDX;
Table altered.
SQL> DROP TABLE products CASCADE CONSTRAINTS PURGE;
Table dropped.
SQL> CREATE TABLE product (product_id INTEGER NOT NULL,
prod_unit_price INTEGER,
description VARCHAR2(150)) TABLESPACE USER_DATA2;
Table created.
SQL> ALTER TABLE product ADD CONSTRAINT product_pk PRIMARY KEY (
product_id ) USING
INDEX TABLESPACE USER_INDX;
Table altered.
SQL> DROP TABLE store CASCADE CONSTRAINTS PURGE;
```

```

Table dropped.
SQL> CREATE TABLE store (storeid INTEGER NOT NULL, manager_id
INTEGER, phoneno
CHAR(15), zip_code CHAR(5)) TABLESPACE USER_DATA2;
Table created.
SQL> ALTER TABLE store ADD CONSTRAINT store_pk PRIMARY KEY ( storeid
) USING INDEX
TABLESPACE USER_INDX;
Table altered.
SQL> DROP TABLE supplier CASCADE CONSTRAINTS PURGE;
Table dropped.
SQL> CREATE TABLE supplier (supplierid INTEGER NOT NULL, suppliername
VARCHAR2(150),
zipcode CHAR(5), product_id INTEGER, phoneno INTEGER, email
VARCHAR2(100), contactname
VARCHAR2(150)) TABLESPACE USER_DATA2;
Table created.
SQL> ALTER TABLE supplier ADD CONSTRAINT supplier_pk PRIMARY KEY (
supplierid ) USING
INDEX TABLESPACE USER_INDX;
Table altered.
SQL> DROP TABLE zip CASCADE CONSTRAINTS PURGE;
Table dropped.
SQL> CREATE TABLE zip (zipcode CHAR(5) NOT NULL, city CHAR(25), state
CHAR(2))
TABLESPACE USER_DATA2;
Table created.
SQL> ALTER TABLE zip ADD CONSTRAINT zip_pk PRIMARY KEY (zipcode)
USING INDEX TABLESPACE
USER_INDX;
Table altered.

```

Foreign Keys

```

SQL> ALTER TABLE customer ADD CONSTRAINT customer_zip_fk FOREIGN KEY
(zip_code)
REFERENCES zip ( zipcode );
Table altered.
SQL> ALTER TABLE employees ADD CONSTRAINT employees_manager_fk
FOREIGN KEY (

```

```
manager_id ) REFERENCES manager ( managerid );
Table altered.
SQL> ALTER TABLE employees ADD CONSTRAINT employees_store_fk FOREIGN
KEY ( store_id )
REFERENCES store ( storeid );
Table altered.
SQL> ALTER TABLE manager ADD CONSTRAINT manager_zip_fk FOREIGN KEY (
zip_code )
REFERENCES zip ( zipcode );
Table altered.
SQL> ALTER TABLE order_item ADD CONSTRAINT order_item_orders_fk
FOREIGN KEY (
orderid ) REFERENCES orders ( order_id );
Table altered.
SQL> ALTER TABLE order_item ADD CONSTRAINT order_item_product_fk
FOREIGN KEY (
productid ) REFERENCES product ( product_id );
Table altered.
SQL> ALTER TABLE orders ADD CONSTRAINT orders_customer_fk FOREIGN KEY (
customerid ) REFERENCES customer ( customer_id );
Table altered.
SQL> ALTER TABLE employees ADD CONSTRAINT employees_orders_fk FOREIGN
KEY ( order_id )
REFERENCES orders (order_id);
Table altered.
SQL> ALTER TABLE payment ADD CONSTRAINT payment_orders_fk FOREIGN KEY
( orderid )
REFERENCES orders ( order_id );
Table altered.
SQL> ALTER TABLE supplier ADD CONSTRAINT supplier_product_fk FOREIGN
KEY ( product_id )
REFERENCES product ( product_id );
Table altered.
SQL> ALTER TABLE store ADD CONSTRAINT store_manager_fk FOREIGN KEY (
manager_id )
REFERENCES manager ( managerid );
Table altered.
SQL> ALTER TABLE store ADD CONSTRAINT store_zip_fk FOREIGN KEY (
```

```

zip_code ) REFERENCES
zip ( zipcode );
Table altered.

SQL> ALTER TABLE supplier ADD CONSTRAINT supplier_zip_fk FOREIGN KEY
( zipcode )
REFERENCES zip ( zipcode );
Table altered.

```

Creating Partitioned Table

```

SQL> drop table ZIP cascade constraints purge;
Table dropped.

SQL> create table ZIP (ZIPCODE char(5) NOT NULL,CITY char(25),STATE
char(2),
constraint pk_zip primary key (ZIPCODE) using index storage (initial
25k next 12k)
tablespace USER_INDX) partition by range (state) (partition zip_part1
values less than ('N')
pctfree 5 pctused 90 tablespace USER_DATA1 storage (initial 1500k
next 200k pctincrease 0 maxextents 5), partition zip_part2 values
less than (MAXVALUE) pctfree 5 pctused 90
tablespace USER_DATA2 storage (initial 1500k next 200k pctincrease 0
maxextents 5));
Table created.

SQL> drop table CUSTOMER cascade constraints purge;
Table dropped.

SQL> create table CUSTOMER (ACCOUNT_NUMBER number(10) NOT NULL, SSN
number(9) NOT
NULL,FIRST_NAME char(10), MIDDLE_INITL char(1), LAST_NAME char(20)
NOT NULL,FIRST_LINE_AD
char(20), FK_ZIPCODE char(5), constraint pk_customer primary key
(ACCOUNT_NUMBER) using
index storage (initial 25k next 12k) tablespace USER_INDX, constraint
fk_zip foreign key
(FK_ZIPCODE) references ZIP(ZIPCODE)) pctfree 0 pctused 90 storage
(initial 50k next 10k
pctincrease 0 maxextents 5) tablespace USER_DATA1;
Table created.

SQL> create index cust_zip_idx on CUSTOMER(FK_ZIPCODE)
storage(initial 10k next 10k)

```

```
tablespace USER_INDX;
Index created.
```

Synonyms are created

```
CREATE OR REPLACE PUBLIC SYNONYM CUSTOMER FOR HANIF.CUSTOMER;
CREATE OR REPLACE PUBLIC SYNONYM PRODUCT FOR HANIF.PRODUCT;
CREATE OR REPLACE PUBLIC SYNONYM ORDERS FOR HANIF.ORDERS;
CREATE OR REPLACE PUBLIC SYNONYM ORDER_ITEM FOR HANIF.ORDER_ITEM;
CREATE OR REPLACE PUBLIC SYNONYM ZIP FOR HANIF.ZIP;
CREATE OR REPLACE PUBLIC SYNONYM PAYMENT FOR HANIF.PAYMENT;
CREATE OR REPLACE PUBLIC SYNONYM SUPPLIER FOR HANIF.SUPPLIER;
CREATE OR REPLACE PUBLIC SYNONYM EMPLOYEES FOR HANIF.EMPLOYEES;
CREATE OR REPLACE PUBLIC SYNONYM STORE FOR HANIF.STORE;
CREATE OR REPLACE PUBLIC SYNONYM MANAGER FOR HANIF.MANAGER;
```

To take notes of the users and their activity, the DBA will show what a user is executing and their details:

```
COL USERNAME FOR A10
COL MACHINE FOR A10
COL SCHEMANAME FOR A10
COL OSUSER FOR A12
COL PROGRAM FOR A30
COL MODULE FOR A20
COL EVENT FOR A30
COL STATUS FOR A8
COL SERVICE_NAME FOR A12
COL STATE FOR A8
SQL> SELECT USERNAME, STATUS,SCHEMANAME,OSUSER,MACHINE, PROGRAM,
MODULE, LOGON_TIME, LAST_CALL_ET,
STATE, SERVICE_NAME, EVENT, SQL_ID FROM V$SESSION WHERE
USERNAME= '&USERNAME';
Enter value for username: HANIF
USERNAME STATUS SCHEMANAME OSUSER MACHINE PROGRAM MODULE
LOGON_TIM LAST_CALL_ET STATE SERVICE_NAME EVENT SQL_ID
-----
-----
```

```
HANIF INACTIVE HANIF oracle NIXORA01 sqlplus@NIXORA01 (TNS V1-V3)
SQL*Plus
30-OCT-22 0 WAITING SYS$USERS SQL*Net message from client
3ws0hu2bdjgt0
1 row selected.
SQL>
```

The event column will inform the DBA or system administrator what the user is doing.

The “db file sequential read” entails reading from a block sequentially and is an I/O operation.

“SQL*net message from client” means waiting for the next query to be fired and to complete its execution. SQL_ID will allow the user to know what other users are doing in terms of SQL execution.

Export and Import

The database administrator user creates test table to test the export and import function to demonstrate capability.

```

Version 19.3.0.0.0
Copyright (c) 1982, 2019, Oracle and/or its affiliates. All rights reserved.

Username: hanifdba
Password:

Connected to: Oracle Database 19c Enterprise Edition Release 19.0.0.0.0 - Production
Version 19.3.0.0.0
Export done in US7ASCII character set and AL16UTF16 NCHAR character set
server uses AL32UTF8 character set (possible charset conversion)

About to export specified users ...
. exporting pre-schema procedural objects and actions
. exporting foreign function library names for user TESTUSER
. exporting PUBLIC type synonyms
. exporting private type synonyms
. exporting object type definitions for user TESTUSER
About to export TESTUSER's objects ...
. exporting database links
. exporting sequence numbers
. exporting cluster definitions
. about to export TESTUSER's tables via Conventional Path ...
. exporting synonyms
. exporting views
. exporting stored procedures
. exporting operators
. exporting referential integrity constraints
. exporting triggers
. exporting indextypes
. exporting bitmap, functional and extensible indexes
. exporting posttables actions
. exporting materialized views
. exporting snapshot logs
. exporting job queues
. exporting refresh groups and children
. exporting dimensions
. exporting post-schema procedural objects and actions
. exporting statistics
Export terminated successfully without warnings.
[oracle@NIXORA01 StudentFirst]$ 

```

RMAN is most efficient for backup and recovery to the database though it cannot backup a single table or schema like the export-import functions can. For schema level the DBA will use export-import, otherwise RMAN should be utilized.

Loading .ctl and .dat files using SQL*Loader

The SQL*Loader utility will load the respective .ctl and .dat files for the order tables. SQL*Loader acts as a utility for the .ctl file parameters to be pass through resulting in said utility

opening the .ctl file and loading the data. Store, manager, employee, and supplier table data is randomly generate to produce a sample data using the Mockaroo website (Moomin, 2022). Data loading is limited to 1000 tables max and 50000 max according to the load files in the index. The loading is done through Linux, fetching the tables in the table directory [oracle@NIXORA01 dir_6tables]. The process is shown below:

```

sqlldr userid=hanif/oracle
control=/u01/app/oracle/repository/dir_6tables/zip.ctl
data=/u01/app/oracle/repository/dir_6tables/hb_zip.dat
bad=/u01/app/oracle/repository/dir_6tables/zip.bad
discard=/u01/app/oracle/repository/dir_6tables/zip.dsc
log=/u01/app/oracle/repository/dir_6tables/zip.log direct=y
SQL*Loader: Release 19.0.0.0.0 - Production on Wed Oct 12 19:43:28
2022 Version 19.3.0.0.
Path used: Direct, LOAD=50000
Load completed - logical record count 50000.
Table ZIP:
      50000 Rows successfully loaded.
Check the log file:
      /u01/app/oracle/repository/dir_6tables/store.log
for more information about the load.
[oracle@NIXORA01 dir_6 tables]$
sqlldr userid=hanif/oracle
control=/u01/app/oracle/repository/dir_6tables/orders.ctl
data=/u01/app/oracle/repository/dir_6tables/order.dat
bad=/u01/app/oracle/repository/dir_6tables/orders.bad
discard=/u01/app/oracle/repository/dir_6tables/orders.dsc
log=/u01/app/oracle/repository/dir_6tables/orders.log direct=y
SQL*Loader: Release 19.0.0.0.0 - Production on Wed Oct 12 19:43:28
2022 Version 19.3.0.0.
Path used: Direct, LOAD=1000
Load completed - logical record count 1000.
Table ORDERS:
      1000 Rows successfully loaded.
Check the log file:
      /u01/app/oracle/repository/dir_6tables/store.log

```

```
for more information about the load.  
[oracle@NIXORA01 dir_6 tables]$  
sqlldr userid=hanif/oracle  
control=/u01/app/oracle/repository/dir_6tables/product.ctl  
data=/u01/app/oracle/repository/dir_6tables/product.dat  
bad=/u01/app/oracle/repository/dir_6tables/product.bad  
discard=/u01/app/oracle/repository/dir_6tables/product.dsc  
log=/u01/app/oracle/repository/dir_6tables/product.log direct=y  
SQL*Loader: Release 19.0.0.0.0 - Production on Wed Oct 12 19:43:28  
2022 Version 19.3.0.0.  
Path used: Direct, LOAD=1000  
Load completed - logical record count 1000.  
Table PRODUCT:  
      1000 Rows successfully loaded.  
Check the log file:  
      /u01/app/oracle/repository/dir_6tables/store.log  
for more information about the load.  
[oracle@NIXORA01 dir_6 tables]$  
sqlldr userid=hanif/oracle  
control=/u01/app/oracle/repository/dir_6tables/order_item.ctl  
data=/u01/app/oracle/repository/dir_6tables/order_item.dat  
bad=/u01/app/oracle/repository/dir_6tables/order_item.bad  
discard=/u01/app/oracle/repository/dir_6tables/order_item.dsc  
log=/u01/app/oracle/repository/dir_6tables/order_item.log direct=y  
SQL*Loader: Release 19.0.0.0.0 - Production onWed Oct 12 19:43:28  
2022 Version 19.3.0.0.  
Path used: Direct, LOAD=1000  
Load completed - logical record count 1000.  
Table ORDER_ITEM:  
      1000 Rows successfully loaded.  
Check the log file:  
      /u01/app/oracle/repository/dir_6tables/order_item.log  
for more information about the load.  
[oracle@NIXORA01 dir_6 tables]$  
sqlldr userid=hanif/oracle  
control=/u01/app/oracle/repository/dir_6tables/payment.ctl  
data=/u01/app/oracle/repository/dir_6tables/payment.dat  
bad=/u01/app/oracle/repository/dir_6tables/payment.bad  
discard=/u01/app/oracle/repository/dir_6tables/payment.dsc
```

```
log=/u01/app/oracle/repository/dir_6tables/payment.log direct=y
SQL*Loader: Release 19.0.0.0.0 - Production onWed Oct 12 19:43:28
2022 Version 19.3.0.0.
Path used: Direct, LOAD=1000
Load completed - logical record count 1000.
Table PAYMENT:
    1000 Rows successfully loaded.
Check the log file:
    /u01/app/oracle/repository/dir_6tables/store.log
for more information about the load.
[oracle@NIXORA01 dir_6 tables]$
sqlldr userid=hanif/oracle
control=/u01/app/oracle/repository/dir_6tables/customer.ctl
data=/u01/app/oracle/repository/dir_6tables/customer.dat
bad=/u01/app/oracle/repository/dir_6tables/customer.bad
discard=/u01/app/oracle/repository/dir_6tables/customer.dsc
log=/u01/app/oracle/repository/dir_6tables/customer.log direct=y
SQL*Loader: Release 19.0.0.0.0 - Production onWed Oct 12 19:43:28
2022 Version 19.3.0.0.
Path used: Direct, LOAD=1000
Load completed - logical record count 1000.
Table CUSTOMER:
    1000 Rows successfully loaded.
Check the log file:
    /u01/app/oracle/repository/dir_6tables/store.log
for more information about the load.
[oracle@NIXORA01 dir_6 tables]$
sqlldr userid=hanif/oracle
control=/u01/app/oracle/repository/dir_6tables/employee.ctl
data=/u01/app/oracle/repository/dir_6tables/employee.dat
bad=/u01/app/oracle/repository/dir_6tables/employee.bad
discard=/u01/app/oracle/repository/dir_6tables/employee.dsc
log=/u01/app/oracle/repository/dir_6tables/employee.log direct=y
SQL*Loader: Release 19.0.0.0.0 - Production on Wed Oct 12 19:43:28
2022 Version 19.3.0.0.
Path used: Direct, LOAD=1000
Load completed - logical record count 1000.
Table EMPLOYEE:
    1000 Rows successfully loaded.
```

Check the log file:

```
/u01/app/oracle/repository/dir_6tables/store.log  
for more information about the load.  
[oracle@NIXORA01 dir_6 tables]$  
sqlldr userid=hanif/oracle  
control=/u01/app/oracle/repository/dir_6tables/manager.ctl  
data=/u01/app/oracle/repository/dir_6tables/manager.dat  
bad=/u01/app/oracle/repository/dir_6tables/manager.bad  
discard=/u01/app/oracle/repository/dir_6tables/manager.dsc  
log=/u01/app/oracle/repository/dir_6tables/manager.log direct=y  
SQL*Loader: Release 19.0.0.0.0 - Production onWed Oct 12 19:43:28  
2022 Version 19.3.0.0.  
Path used: Direct, LOAD=1000  
Load completed - logical record count 1000.
```

Table MANAGER:

```
1000 Rows successfully loaded.
```

Check the log file:

```
/u01/app/oracle/repository/dir_6tables/store.log  
for more information about the load.  
[oracle@NIXORA01 dir_6 tables]$  
sqlldr userid=hanif/oracle  
control=/u01/app/oracle/repository/dir_6tables/store.ctl  
data=/u01/app/oracle/repository/dir_6tables/store.dat  
bad=/u01/app/oracle/repository/dir_6tables/store.bad  
discard=/u01/app/oracle/repository/dir_6tables/store.dsc  
log=/u01/app/oracle/repository/dir_6tables/store.log direct=y  
SQL*Loader: Release 19.0.0.0.0 - Production onWed Oct 12 19:43:28  
2022 Version 19.3.0.0.  
Path used: Direct, LOAD=1000  
Load completed - logical record count 1000.
```

Table STORE:

```
1000 Rows successfully loaded.
```

Check the log file:

```
/u01/app/oracle/repository/dir_6tables/store.log  
for more information about the load.  
[oracle@NIXORA01 dir_6 tables]$  
sqlldr userid=hanif/oracle  
control=/u01/app/oracle/repository/dir_6tables/supplier.ctl  
data=/u01/app/oracle/repository/dir_6tables/supplier.dat
```

```

bad=/u01/app/oracle/repository/dir_6tables/supplier.bad
discard=/u01/app/oracle/repository/dir_6tables/supplier.dsc
log=/u01/app/oracle/repository/dir_6tables/supplier.log direct=y
SQL*Loader: Release 19.0.0.0.0 - Production on Wed Oct 12 19:43:28
2022 Version 19.3.0.0.
Path used: Direct, LOAD=1000
Load completed - logical record count 1000.
Table SUPPLIER:
      998 Rows successfully loaded.
Check the log file:
      /u01/app/oracle/repository/dir_6tables/store.log
for more information about the load.
[oracle@NIXORA01 dir_6 tables]$

```

Table Analyze

The user will run clauses to analyze the tables. ANALYZE TABLE command is used along with the DBMS_STATS including verifying that objects are valid. First, the tables are analyzed so that oracle can collect statistics.

```

SQL> analyze table ZIP compute statistics;
Table analyzed.
SQL> analyze table ORDER_ITEM compute statistics;
Table analyzed.
SQL> analyze table ORDERS compute statistics;
Table analyzed.
SQL> analyze table PAYMENT compute statistics;
Table analyzed.
SQL> analyze table PRODUCT compute statistics;
Table analyzed.
SQL> analyze table CUSTOMER compute statistics;
Table analyzed.
SQL> analyze table EMPLOYEES compute statistics;
Table analyzed.
SQL> analyze table SUPPLIER compute statistics;
Table analyzed.
SQL> analyze table STORE compute statistics;
Table analyzed.

```

```
SQL> analyze table MANAGER compute statistics;
Table analyzed.
```

The DBA can then use the LAST_ANALYZE, STALE_STATS and NUM_ROWS clauses as examples of user statistics.

```
SQL>SET LINES 350
SQL>SET PAGES 5000
SQL>col table_name for a25
SQL>select table_name, last_analyzed, stale_stats from
user_tab_statistics where
table_name in ('ZIP', 'EMPLOYEES', 'SUPPLIER', 'ORDER_ITEM',
'ORDERS', 'PAYMENT', 'PRODUCT', 'STORE', 'CUSTOMER', 'MANAGER')
ORDER BY 2 DESC;
```

Result:

```
SQL> SET PAGES 5000
SQL> SET LINES 350
SQL> COL TABLE_NAME FOR A25
SQL> SELECT TABLE_NAME, LAST_ANALYZED, STALE_STATS, NUM_ROWS FROM USER_TAB_STATISTICS WHERE TABLE_NAME IN ('ZIP', 'EMPLOYEES', 'SUPPLIER', 'ORDER_ITEM', 'ORDERS', 'PAYMENT', 'PRODUCT', 'STORE', 'CUSTOMER', 'MANAGER') ORDER BY 2 DESC;

TABLE_NAME          LAST_ANAL STALE_S    NUM_ROWS
-----  -----  -----
ZIP                20-OCT-22 NO        50000
ZIP                20-OCT-22 NO        24377
ZIP                20-OCT-22 NO        25623
CUSTOMER           12-OCT-22 NO        1000
PAYMENT            12-OCT-22 NO        1000
ORDER_ITEM         12-OCT-22 NO        1000
ORDERS             12-OCT-22 NO        1000
PRODUCT            12-OCT-22 NO        1000
MANAGER            12-OCT-22 NO        1000
STORE              12-OCT-22 NO        1000
SUPPLIER           12-OCT-22 NO        998
EMPLOYEES          12-OCT-22 NO        1000

12 rows selected.

SQL> █
```

The statistics that were gathered from the database is done using “DBMS_STATS”

```
SQL>
BEGIN
```

```
DBMS_STATS.GATHER_TABLE_STATS
(
  OWNNAME =>'HANIF', TABNAME => 'ZIP', ESTIMATE_PERCENT =>
  DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
  AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
  DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.

SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
  OWNNAME =>'HANIF', TABNAME => 'PAYMENT', ESTIMATE_PERCENT =>
  DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
  AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
  DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.

SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
  OWNNAME =>'HANIF', TABNAME => 'ORDERS', ESTIMATE_PERCENT =>
  DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
  AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
  DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.

SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
  OWNNAME =>'HANIF', TABNAME => 'PRODUCT', ESTIMATE_PERCENT =>
  DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
  AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
  DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.
```

```
SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
OWNNAME =>'HANIF', TABNAME => 'ORDER_ITEM', ESTIMATE_PERCENT =>
DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.
SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
OWNNAME =>'HANIF', TABNAME => 'CUSTOMER', ESTIMATE_PERCENT =>
DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.
SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
OWNNAME =>'HANIF', TABNAME => 'EMPLOYEE', ESTIMATE_PERCENT =>
DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.
SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
OWNNAME =>'HANIF', TABNAME => 'MANAGER', ESTIMATE_PERCENT =>
DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
DEGREE => 2);
END;
/
```

```

PL/SQL procedure successfully completed.
SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
OWNNAME =>'HANIF', TABNAME => 'STORE', ESTIMATE_PERCENT =>
DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.
SQL>
DBMS_STATS.GATHER_TABLE_STATS
(
OWNNAME =>'HANIF', TABNAME => 'SUPPLIER', ESTIMATE_PERCENT =>
DBMS_STATS.AUTO_SAMPLE_SIZE, METHOD_OPT => 'FOR ALL COLUMNS SIZE
AUTO', CASCADE => TRUE, GRANULARITY => 'ALL', NO_INVALIDATE => FALSE,
DEGREE => 2);
END;
/
PL/SQL procedure successfully completed.

```

Table Sizes

The table sizes are calculated and compared with the capacity planning. The product table capacity is 384 kb to the 1990 kb. This can be attributed to the only 1000 rows that comprise the product table over the 14,000 estimated. Order_item, customer, product and orders are among the calculations that were undervalued. The rest are overcalculated. There isn't substantial accuracy based on assumptions, but it's a good practice to undergo capacity planning for the database project plan. Project managers must be prepared for over or underestimates. Below is the actual table size:

```

SQL> select SEGMENT_NAME as "OBJECT_NAME", SEGMENT_TYPE as
"OBJECT_TYPE", sum(bytes)/(1024) as "TABLE_SIZE(KB)" from
dba_segments where segment_name IN('EMPLOYEES', 'SUPPLIER',
'ORDER_ITEM', 'ORDERS', 'PAYMENT', 'PRODUCT', 'STORE', 'CUSTOMER',

```

```
'MANAGER', 'ZIP') group by SEGMENT_NAME, SEGMENT_TYPE;
OBJECT OBJECT
NAME          TYPE      TABLE_SIZE(KB)
-----
ZIP TABLE      PARTITION 5120
PAYMENT        TABLE    64
STORE          TABLE    64
SUPPLIER        TABLE   128
PRODUCT        TABLE   384
ORDER_ITEM      TABLE    64
EMPLOYEES      TABLE   128
CUSTOMER        TABLE   128
ORDERS          TABLE    64
MANAGER         TABLE    64
10 rows selected.
```

Monitoring the Database Statistics

V\$SYSSTAT is used to provide access to various database statistics collected. It is important to note that numbers aren't guaranteed to remain constant from one reading to another, so the DBA should take into consideration the date and time these statistic monitoring sessions occur. The DBA must log in to the SYS user to monitor the database. Physical, cache, and

```
SQL> select name,value from v$sysstat where name like 'physical%'
order by value desc;
NAME
VALUE
-----
physical read total bytes
765639680
physical read bytes
191111168
physical write total bytes
121629237
physical write bytes 24207360
physical read total IO requests
42962
```

```
physical reads 23329
physical reads cache 13355
physical reads direct (lob) 9973
physical reads direct 9973
physical read IO requests
8268
physical write total IO requests
7984
physical reads cache prefetch 5383
physical writes
2955
physical writes from cache
2792
physical writes non checkpoint
2075
physical write IO requests
1386
physical read total multi block requests 215
physical writes direct (lob) 163
physical writes direct
163
physical write total multi block requests 52
physical reads prefetch warmup
6
physical reads direct for securefile flashback block new 0
physical reads cache for securefile flashback block new 0
physical reads for flashback new
0
physical reads retry corrupt 0
physical writes direct temporary tablespace 0
physical reads direct temporary tablespace 0
physical reads for data transfer
0
physical write snap IO requests new allocations 0
physical read snap bytes copy 0
physical read snap bytes base 0
physical read snap IO requests no data
0
physical read snap IO requests copy 0
```

```

physical read snap IO requests base 0
physical write total bytes optimized 0
physical write requests optimized 0
physical read partial requests
0
physical read total bytes optimized 0
physical read requests optimized
0
physical read flash cache hits
0
40 rows selected.
Elapsed: 00:00:00.01

```

```

SQL> select name,value from v$sysstat where name like '%cache%' order
by value desc;
NAME VALUE
-----
-----
logical read bytes from cache 5093851136
consistent gets from cache 555010
session cursor cache hits 102638
db block gets from cache 66798
db block gets from cache (fastpath) 54339
physical reads cache 16248
physical reads cache prefetch 7121
physical writes from cache 5766
session cursor cache count 2582
LOB table id lookup cache misses 3
Commit SCN cached 1
cell writes to flash cache 0
cell overwrites in flash cache 0
cell partial writes in flash cache 0
cell writes to flash cache for temp IO 0
physical read flash cache hits 0
flash cache inserts 0
flash cache eviction: invalidated 0
flash cache eviction: buffer pinned 0
flash cache eviction: aged out 0

```

```
flash cache insert skip: not current 0
flash cache insert skip: DBWR overloaded 0
flash cache insert skip: exists 0
flash cache insert skip: not useful 0
flash cache insert skip: modification 0
flash cache insert skip: corrupt 0
recovery block gets from cache 0
flashback cache read optimizations for block new 0
flashback securefile cache read optimizations for block new 0
physical reads cache for securefile flashback block new 0
gc cluster flash cache reads served 0
gc cluster flash cache reads received 0
gc cluster flash cache received read time 0
gc cluster flash cache reads failure 0
gc flash cache reads served 0
gc flash cache served read time 0
BA count - cache miss 0
BA count - cache denied 0
cell physical IO bytes saved by columnar cache 0
cell smart IO session cache lookups 0
cell smart IO session cache hits 0
cell smart IO session cache soft misses 0
cell smart IO session cache hard misses 0
cell smart IO session cache hwm 0
txn cache local read hits 0
txn cache local read misses 0
txn cache local loscn read misses 0
txn cache remote loscn read misses 0
txn cache local writes 0
txn cache local usn hash write misses 0
txn cache local usn hash read misses 0
txn cache remote usn hash read misses 0
txn cache remote read hits 0
txn cache remote read misses 0
txn cache remote writes 0
txn cache remote active read misses 0
txn cache local ac read misses 0
txn cache redo sync reads 0
txn cache remote copy hits 0
```

```
txn cache remote copy misses 0
txn cache remote ts/inst mismatch 0
txn cache remote proc misses 0
txn cache local proc misses 0
txn cache upper bound misses 0
txn cache local preset misses 0
txn cache local seq num misses 0
txn cache remote seq num misses 0
txn cache local empty slot misses 0
txn cache remote empty slot misses 0
txn cache local sync commit scn misses 0
txn cache remote sync commit scn misses 0
txn cache remote read msg sent 0
txn cache remote requested xid count 0
txn cache remote fetch double pass 0
txn cache lookup usn exceeds max 0
txn cache lookup slot exceeds max 0
txn cache lookup fail read even version 0
txn cache lookup conflicting write exact 0
txn cache lookup conflicting write upper 0
txn cache init usn slot mismatch 0
txn cache init usn exceeds max 0
txn cache write usn exceeds max 0
txn cache write slot mismatch 0
IM populate cache read time (ms) 0
IM scan CUs CR cache hit 0
IM scan CUs CR cache miss 0
table scans (cache partitions) 0
cell blocks processed by cache layer 0
cell commit cache queries 0
cell transactions found in commit cache 0
cell blocks helped by commit cache 0
securefile dedup wapp cache miss 0
IM scan rows cache delta exists 0
IM scan rows cache populate 0
IM scan rows cache no delta 0
IM scan blocks cache 0
key vector CUs processed using cached integer join keys 0
cell flash cache read hits 0
```

```

cell flash cache read hits for controlfile reads 0
cell flash cache read hits for smart IO 0
cell flash cache read hits for temp IO 0
cell ram cache read hits 0
cell pmem cache read hits 0
cell pmem cache writes 0
104 rows selected.
Elapsed: 00:00:00.01
SQL>

```

Statistic Interpretation

Some statistics stick out: the CPU used by the session, sorts, session pga memory, file i/o wait time, user i/o wait time, redo size, redo wastage, enqueue timeouts and waits, db block gets and from the cache as well. Five of these statistics are further assessed: CPU, sorts, user i/o wait time, redo wastage, and enqueue timeouts.

```

SQL> select name,value from v$sysstat where name in ('CPU used by
this session', 'sorts
(memory)', 'user I/O wait time', 'redo wastage','enqueue timeouts')
order by value desc;
NAME                                VALUE
-----                                 -----
redo wastage                         37124
sorts (memory)                      8324
CPU used by this session            1764
user I/O wait time                 1360
enqueue timeouts                     5
5 rows selected.
Elapsed: 00:00:00.01

```

To test the effects of workload, the workload query that runs for 1:00:00 is ran and its effect on the statistics are noted.

```

SQL> ed
Wrote file afiedt.buf
1 DECLARE
2 P_CUSTOMER_ID HANIF.CUSTOMER.CUSTOMER_ID%TYPE;

```

```

3 P_FIRSTNAME HANIF.CUSTOMER.FIRSTNAME%TYPE;
4 P_LASTNAME HANIF.CUSTOMER.LASTNAME%TYPE;
5 P_ADDRESS HANIF.CUSTOMER.ADDRESS%TYPE;
6 P_ORDER_ID HANIF.ORDERS.ORDER_ID%TYPE;
7 P_ORDER_METHOD HANIF.ORDERS.ORDER_METHOD%TYPE;
8 P_CUSTOMERID HANIF.ORDERS.CUSTOMERID%TYPE;
9 P_PRODUCT_ID HANIF.PRODUCT.PRODUCT_ID%TYPE;
10 P_DESCRIPTION HANIF.PRODUCT.DESCRIPTION%TYPE;
11 P_QUANTITY HANIF.ORDER_ITEM.QUANTITY%TYPE;
12 i NUMBER;
13 BEGIN
14 i :=1;
15 FOR i IN 1..100000000 LOOP
16 SELECT C.CUSTOMER_ID, C.FIRSTNAME, C.LASTNAME, C.ADDRESS,
O.ORDER_ID, O.ORDER_METHOD
INTO P_CUSTOMER_ID, P_FIRSTNAME, P_LASTNAME, P_ADDRESS, P_ORDER_ID,
P_ORDER_METHOD FROM
HANIF.CUSTOMER C INNER JOIN HANIF.ORDERS O ON C.CUSTOMER_ID =
O.CUSTOMERID AND ROWNUM=1;
17 SELECT O.ORDER_ID, O.CUSTOMERID, O.ORDER_METHOD, P.PRODUCT_ID,
P.DESCRIPTION,
OT.QUANTITY INTO P_ORDER_ID, P_CUSTOMERID, P_ORDER_METHOD,
P_PRODUCT_ID, P_DESCRIPTION,
P_QUANTITY FROM HANIF.PRODUCT P INNER JOIN HANIF.ORDER_ITEM OT ON
P.PRODUCT_ID =
OT.PRODUCTID INNER JOIN HANIF.ORDERS O ON OT.ORDERID = O.ORDER_ID AND
ROWNUM=1;
18 END LOOP;
19* END;
SQL> /
PL/SQL procedure successfully completed.
Elapsed: 01:01:27.78

```

The statistics collection is queried again:

```

SQL> select name,value from v$sysstat where name in ('CPU used by
this session', 'sorts
(memory)', 'user I/O wait time', 'redo wastage','enqueue timeouts')
order by value desc;
NAME                      VALUE

```

```

-----
CPU used by this session      382931
redo wastage                  177508
sorts (memory)                22709
user I/O wait time            1528
enqueue timeouts               6
5 rows selected.

```

The intensive query had an effect to the degree that the CPU used multiplied by 10.3, redo wastage by 21.56, sorts by 12.87, user i/o wait time increase by 168 and enqueue timeouts increased by 1. Redo wastage benefits sequential i/o pattern but the high redo wastage indicates a very active log writer (Burleson, n.d.b). Sorts memory metric indicates how much the sort is performed completely, no sorts is most beneficial on the workloads (Burleson, n.d.c). CPU used by the session multiplied by 10 times. This is a considerable amount of CPU usage and SQL has optimizer cost tools to counter this (Burleson, n.d.a). The changes reflect what the DBA would expect from a workload of only reads. Reads don't create locks, therefore enqueue changes by one. The database is restarted and the write workload is ran:

82

Note: The database was restarted after Read workload execution.

Statistics collection before write workload execution:

```

SQL> select name,value from v$sysstat where name in ('CPU used by
this session', 'sorts
(memory)', 'user I/O wait time', 'redo wastage','enqueue timeouts')
order by value desc;
NAME                 VALUE
-----
redo wastage          29628
sorts (memory)        16467
CPU used by this session    1247
user I/O wait time      1175

```

```
enqueue timeouts          5
5 rows selected.
Elapsed: 00:00:00.02
```

```
SQL> DROP TABLE HANIF.ORDER_ITEM_1 CASCADE CONSTRAINTS PURGE;
Table dropped.
Elapsed: 00:00:00.60
SQL> CREATE TABLE HANIF.ORDER_ITEM_1 AS SELECT * FROM
HANIF.ORDER_ITEM WHERE 1=2;
Table created.
Elapsed: 00:00:00.22
SQL> TRUNCATE TABLE HANIF.ORDER_ITEM_1;
Table truncated.
Elapsed: 00:00:00.09
SQL> SET SERVEROUTPUT ON
SQL> BEGIN
  INSERT INTO HANIF.ORDER_ITEM_1 SELECT * FROM HANIF.ORDER_ITEM CONNECT
  BY LEVEL <= 2;
  DBMS_OUTPUT.PUT_LINE('ROWS INSERTED = ' || SQL%ROWCOUNT);
  COMMIT;
END;
/
ROWS INSERTED = 1211100
PL/SQL procedure successfully completed.
Elapsed: 00:01:48.10
SYS>
```

The statistics collection is queried again.

```
SQL> select name,value from v$sysstat where name in ('CPU used by
this session', 'sorts
(memory)', 'user I/O wait time', 'redo wastage','enqueue timeouts')
order by value desc;
NAME                  VALUE
-----
redo wastage          204672
sorts (memory)        25777
CPU used by this session    13769
user I/O wait time      1572
```

```
enqueue timeouts      5
5 rows selected.
Elapsed: 00:00:00.03
```

The intensive query had an effect to the degree that the CPU used multiplied by 6.93, redo wastage by 1.57, sorts by 7.8, user i/o wait time increase by 325 and enqueue timeouts increased by 1. User i/o wait time increased two times more than it did with the read workload and enqueue remained the same, outside of that the read workload made for a more intensive query than the write workload. Regardless their values increased. No lock is created so the enqueue timeouts remained static.

Performance Tuning with Automatic Database Diagnostic Monitor (ADDM)

Since Oracle 10g, the Automatic Database Diagnostic Monitor (ADDM) is a utility introduced in Oracle 10g showcasing overview of the most prominent performance bottlenecks while offering suggestions for improvement (Oracle, 2019). This feature is utilized for the order management system:

```
SQL> set long 1000000
SQL> set pagesize 50000
SQL> column get_clob format a80
SQL> select dbms_advisor.get_task_report(task_name, 'TEXT', 'ALL')
  as ADDM_report from dba_advisor_tasks where task_id=(select
max(t.task_id) from dba_advisor_tasks t,dba_advisor_log l where
t.task_id = l.task_id and t.advisor_name='ADDM'and l.status=
'COMPLETED');
ADDM_REPORT
-----
-----
ADDM Report for Task 'ADDM:992067541_1_95'
-----
Analysis Period
-----
AWR snapshot range from 94 to 95.
Time period starts at 27-OCT-22 02.00.04 AM
```

```
Time period ends at 27-OCT-22 02.04.18 AM
Analysis Target
-----
Database 'DBST670' with DB ID 992067541.
Database version 19.0.0.0.0.
ADDM performed an analysis of instance DBST670, numbered 1 and hosted
at
NIXORA01.
Activity During the Analysis Period
-----
Total database time was 0 seconds.
The average number of active sessions was 0.
~~~~~
~~~~~
~~~~~There are no findings to report.
~~~~~
~~~~~
~~~~~
Additional Information
-----
Miscellaneous Information
-----
There was no significant database activity to run the ADDM.
1 row selected.
Elapsed: 00:00:00.13
```

The read workload is ran, and the ADDM is subsequently initialized.

```
ADDM_REPORT
-----
-----
ADDM Report for Task 'ADDM:992067541_1_96'
-----
Analysis Period
-----
AWR snapshot range from 95 to 96.
Time period starts at 27-OCT-22 02.04.18 AM
Time period ends at 27-OCT-22 03.00.11 AM
```

Analysis Target

Database 'DBST670' with DB ID 992067541.
Database version 19.0.0.0.0.
ADDM performed an analysis of instance DBST670, numbered 1 and hosted
at
NIXORA01.

Activity During the Analysis Period

Total database time was 1459 seconds.
The average number of active sessions was .44.

Summary of Findings**Description Active Sessions Recommendations****Percent of Activity**

1 Top SQL Statements .4 | 91.22 5

~~~~~  
~~~~~  
~~~~~  
~~~~~

Findings and Recommendations**Finding 1: Top SQL Statements**

Impact is .4 active sessions, 91.22% of total activity.

SQL statements consuming significant database time were found. These
statements offer a good opportunity for performance improvement.

Recommendation 1: SQL Tuning

Estimated benefit is .18 active sessions, 40.54% of total activity.

Action

Run SQL Tuning Advisor on the SELECT statement with SQL_ID
"8h4k7hum8jgz4".

Related Object

SQL statement with SQL_ID 8h4k7hum8jgz4.

SELECT O.ORDER_ID, O.CUSTOMERID, O.ORDER_METHOD, P.PRODUCT_ID,
P.DESCRIPTION, OT.QUANTITY FROM HANIF.PRODUCT P, HANIF.ORDER_ITEM
OT, HANIF.ORDERS O WHERE ROWNUM=1

Rationale

The SQL spent 100% of its database time on CPU, I/O and Cluster waits.

This part of database time may be improved by the SQL Tuning Advisor.

Rationale

Database time for this SQL was divided as follows: 100% for SQL execution, 0% for parsing, 0% for PL/SQL execution and 0% for Java execution.

Rationale

SQL statement with SQL_ID "8h4k7hum8jgz4" was executed 30000000 times and had an average elapsed time of 0.000017 seconds.

Rationale

Top level calls to execute the PL/SQL statement with SQL_ID "ba9c6wwn4rm1d" are responsible for 63% of the database time spent on the SELECT statement with SQL_ID "8h4k7hum8jgz4".

Related Object

SQL statement with SQL_ID ba9c6wwn4rm1d.

DECLARE

P_CUSTOMER_ID

HANIF.CUSTOMER.CUSTOMER_ID%TYPE;

P_FIRSTNAME

HANIF.CUSTOMER.FIRSTNAME%TYPE;

P_LASTNAME

HANIF.CUSTOMER.LASTNAME%TYPE;

P_ADDRESS

HANIF.CUSTOMER.ADDRESS%TYPE;

P_ORDER_ID

HANIF.ORDERS.ORDER_ID%TYPE;

P_ORDER_METHOD

HANIF.ORDERS.ORDER_METHOD%TYPE;

P_CUSTOMERID

HANIF.ORDERS.CUSTOMERID%TYPE;

P_PRODUCT_ID

HANIF.PRODUCT.PRODUCT_ID%TYPE;

P_DESCRIPTION

HANIF.PRODUCT.DESCRIPTION%TYPE;

P_QUANTITY

HANIF.ORDER_ITEM.QUANTITY%TYPE;

i NUMBER;

```

BEGIN
i :=1;
FOR i IN 1..20000000 LOOP
SELECT
C.CUSTOMER_ID, C.FIRSTNAME, C.LASTNAME, C.ADDRESS, O.ORDER_ID,
O.ORDER_METHOD INTO P_CUSTOMER_ID, P_FIRSTNAME, P_LASTNAME,
P_ADDRESS, P_ORDER_ID, P_ORDER_METHOD FROM HANIF.CUSTOMER C,
HANIF.ORDERS O WHERE ROWNUM=1;
SELECT O.ORDER_ID, O.CUSTOMERID, O.ORDER_METHOD, P.PRODUCT_ID,
P.DESCRIPTION, OT.QUANTITY INTO P_ORDER_ID, P_CUSTOMERID,
P_ORDER_METHOD, P_PRODUCT_ID, P_DESCRIPTION, P_QUANTITY FROM
HANIF.PRODUCT P, HANIF.ORDER_ITEM OT, HANIF.ORDERS O WHERE
ROWNUM=1;
END LOOP;
END;

```

Recommendation 2: SQL Tuning

Estimated benefit is .1 active sessions, 22.97% of total activity.

Action

Run SQL Tuning Advisor on the SELECT statement with SQL_ID "gmqm0a80cd8br".

Related Object

SQL statement with SQL_ID gmqm0a80cd8br.

```
SELECT C.CUSTOMER_ID, C.FIRSTNAME, C.LASTNAME, C.ADDRESS, O.ORDER_ID,
O.ORDER_METHOD FROM HANIF.CUSTOMER C, HANIF.ORDERS O WHERE ROWNUM=1
```

Rationale

The SQL spent 100% of its database time on CPU, I/O and Cluster waits.

This part of database time may be improved by the SQL Tuning Advisor.

Rationale

Database time for this SQL was divided as follows: 100% for SQL execution, 0% for parsing, 0% for PL/SQL execution and 0% for Java execution.

Rationale

SQL statement with SQL_ID "gmqm0a80cd8br" was executed 30000000 times and had an average elapsed time of 0.000009 seconds.

Rationale

Top level calls to execute the PL/SQL statement with SQL_ID "ba9c6wwn4rm1d" are responsible for 67% of the database time spent on

```
the SELECT statement with SQL_ID "gmqm0a80cd8br".  
Related Object  
SQL statement with SQL_ID ba9c6wwn4rm1d.  
DECLARE  
P_CUSTOMER_ID  
HANIF.CUSTOMER.CUSTOMER_ID%TYPE;  
P_FIRSTNAME  
HANIF.CUSTOMER.FIRSTNAME%TYPE;  
P_LASTNAME  
HANIF.CUSTOMER.LASTNAME%TYPE;  
P_ADDRESS  
HANIF.CUSTOMER.ADDRESS%TYPE;  
P_ORDER_ID  
HANIF.ORDERS.ORDER_ID%TYPE;  
P_ORDER_METHOD  
HANIF.ORDERS.ORDER_METHOD%TYPE;  
P_CUSTOMERID  
HANIF.ORDERS.CUSTOMERID%TYPE;  
P_PRODUCT_ID  
HANIF.PRODUCT.PRODUCT_ID%TYPE;  
P_DESCRIPTION  
HANIF.PRODUCT.DESCRIPTION%TYPE;  
P_QUANTITY  
HANIF.ORDER_ITEM.QUANTITY%TYPE;  
i NUMBER;  
BEGIN  
i :=1;  
FOR i IN 1..20000000 LOOP  
SELECT  
C.CUSTOMER_ID, C.FIRSTNAME, C.LASTNAME, C.ADDRESS, O.ORDER_ID,  
O.ORDER_METHOD INTO P_CUSTOMER_ID, P_FIRSTNAME, P_LASTNAME,  
P_ADDRESS, P_ORDER_ID, P_ORDER_METHOD FROM HANIF.CUSTOMER C,  
HANIF.ORDERS O WHERE ROWNUM=1;  
SELECT O.ORDER_ID, O.CUSTOMERID, O.ORDER_METHOD, P.PRODUCT_ID,  
P.DESCRIPTION, OT.QUANTITY INTO P_ORDER_ID, P_CUSTOMERID,  
P_ORDER_METHOD, P_PRODUCT_ID, P_DESCRIPTION, P_QUANTITY FROM  
HANIF.PRODUCT P, HANIF.ORDER_ITEM OT, HANIF.ORDERS O WHERE  
ROWNUM=1;  
END LOOP;
```

```

END;

Recommendation 3: SQL Tuning
Estimated benefit is .06 active sessions, 13.51% of total activity.
-----
Action
Run SQL Tuning Advisor on the SELECT statement with SQL_ID
"bcg8x196y0n41".
Related Object
SQL statement with SQL_ID bcg8x196y0n41.
SELECT C.CUSTOMER_ID, C.FIRSTNAME, C.LASTNAME, C.ADDRESS, O.ORDER_ID,
O.ORDER_METHOD FROM HANIF.CUSTOMER C INNER JOIN HANIF.ORDERS O ON
C.CUSTOMER_ID = O.CUSTOMERID AND ROWNUM=1
Rationale
The SQL spent 100% of its database time on CPU, I/O and Cluster
waits.
This part of database time may be improved by the SQL Tuning Advisor.
Rationale
Database time for this SQL was divided as follows: 100% for SQL
execution, 0% for parsing, 0% for PL/SQL execution and 0% for Java
execution.
Rationale
Top level calls to execute the PL/SQL statement with SQL_ID
"57gx14s4sjsah" are responsible for 100% of the database time spent
on
the SELECT statement with SQL_ID "bcg8x196y0n41".
Related Object
SQL statement with SQL_ID 57gx14s4sjsah.
Recommendation 4: SQL Tuning
Estimated benefit is .03 active sessions, 7.43% of total activity.
-----
Action
Investigate the PL/SQL statement with SQL_ID "ba9c6wwn4rm1d" for
possible performance improvements. You can supplement the information
given here with an ASH report for this SQL_ID.
Related Object
SQL statement with SQL_ID ba9c6wwn4rm1d.
DECLARE
P_CUSTOMER_ID
HANIF.CUSTOMER.CUSTOMER_ID%TYPE;

```

```

P_FIRSTNAME
HANIF.CUSTOMER.FIRSTNAME%TYPE;
P_LASTNAME
HANIF.CUSTOMER.LASTNAME%TYPE;
P_ADDRESS
HANIF.CUSTOMER.ADDRESS%TYPE;
P_ORDER_ID
HANIF.ORDERS.ORDER_ID%TYPE;
101
P_ORDER_METHOD
HANIF.ORDERS.ORDER_METHOD%TYPE;
P_CUSTOMERID
HANIF.ORDERS.CUSTOMERID%TYPE;
P_PRODUCT_ID
HANIF.PRODUCT.PRODUCT_ID%TYPE;
P_DESCRIPTION
HANIF.PRODUCT.DESCRIPTION%TYPE;
P_QUANTITY
HANIF.ORDER_ITEM.QUANTITY%TYPE;
i NUMBER;
BEGIN
i :=1;
FOR i IN 1..20000000 LOOP
SELECT
C.CUSTOMER_ID, C.FIRSTNAME, C.LASTNAME, C.ADDRESS, O.ORDER_ID,
O.ORDER_METHOD INTO P_CUSTOMER_ID, P_FIRSTNAME, P_LASTNAME,
P_ADDRESS, P_ORDER_ID, P_ORDER_METHOD FROM HANIF.CUSTOMER C,
HANIF.ORDERS O WHERE ROWNUM=1;
SELECT O.ORDER_ID, O.CUSTOMERID, O.ORDER_METHOD, P.PRODUCT_ID,
P.DESCRIPTION, OT.QUANTITY INTO P_ORDER_ID, P_CUSTOMERID,
P_ORDER_METHOD, P_PRODUCT_ID, P_DESCRIPTION, P_QUANTITY FROM
HANIF.PRODUCT P, HANIF.ORDER_ITEM OT, HANIF.ORDERS O WHERE
ROWNUM=1;
END LOOP;
END;
Rationale
The SQL Tuning Advisor cannot operate on PL/SQL statements.
Rationale
Database time for this SQL was divided as follows: 37% for SQL

```

execution, 0% for parsing, 63% for PL/SQL execution and 0% for Java execution.

Rationale

SQL statement with SQL_ID "ba9c6wwn4rm1d" was executed 1 times and had

an average elapsed time of 755 seconds.

Recommendation 5: SQL Tuning

Estimated benefit is .03 active sessions, 6.76% of total activity.

Action

Run SQL Tuning Advisor on the SELECT statement with SQL_ID "3nmpa03ayt7av".

Related Object

SQL statement with SQL_ID 3nmpa03ayt7av.

```
SELECT O.ORDER_ID, O.CUSTOMERID, O.ORDER_METHOD, P.PRODUCT_ID,
P.DESCRIPTION, OT.QUANTITY FROM HANIF.PRODUCT P INNER JOIN
HANIF.ORDER_ITEM OT ON P.PRODUCT_ID = OT.PRODUCTID INNER JOIN
HANIF.ORDERS O ON OT.ORDERID = O.ORDER_ID AND ROWNUM=1
```

Rationale

The SQL spent 100% of its database time on CPU, I/O and Cluster waits.

This part of database time may be improved by the SQL Tuning Advisor.

Rationale

Database time for this SQL was divided as follows: 100% for SQL execution, 0% for parsing, 0% for PL/SQL execution and 0% for Java execution.

Rationale

Top level calls to execute the PL/SQL statement with SQL_ID "57gx14s4sjsah" are responsible for 100% of the database time spent on

the SELECT statement with SQL_ID "3nmpa03ayt7av".

Related Object

SQL statement with SQL_ID 57gx14s4sjsah.

Additional Information

```
Miscellaneous Information
-----
Wait class "Application" was not consuming significant database time.
Wait class "Commit" was not consuming significant database time.
Wait class "Concurrency" was not consuming significant database time.
Wait class "Configuration" was not consuming significant database
time.
CPU was not a bottleneck for the instance.
Wait class "Network" was not consuming significant database time.
Wait class "User I/O" was not consuming significant database time.
Session connect and disconnect calls were not consuming significant
database
time.
Hard parsing of SQL statements was not consuming significant database
time.
1 row selected.
Elapsed: 00:00:00.16
```

The SQL statements consumed a significant amount of database time and CPU. SQL tuning is recommending to tune the read query causing such a strain. For comparative studies, the write workload is ran. The ADDM report is below.

```
ADDM_REPORT
-----
-----
ADDM Report for Task 'ADDM:992067541_1_101'
-----
Analysis Period
-----
AWR snapshot range from 100 to 101.
Time period starts at 27-OCT-22 10.06.06 AM
Time period ends at 27-OCT-22 10.18.27 AM
Analysis Target
-----
Database 'DBST670' with DB ID 992067541.
Database version 19.0.0.0.0.
ADDM performed an analysis of instance DBST670, numbered 1 and hosted
at
```

NIXORA01.

Activity During the Analysis Period

Total database time was 303 seconds.

The average number of active sessions was .41.

Summary of Findings

Description	Active Sessions	Recommendations
-------------	-----------------	-----------------

Percent of Activity		
---------------------	--	--

-		
1 Top SQL Statements	.41 100	1

2 Session Connect and Disconnect	.01 2.6	1
----------------------------------	-----------	---

3 Hard Parse	.01 2.35	0
--------------	------------	---

Findings and Recommendations

Finding 1: Top SQL Statements

Impact is .41 active sessions, 100% of total activity.

SQL statements consuming significant database time were found. These statements offer a good opportunity for performance improvement.

Recommendation 1: SQL Tuning

Estimated benefit is .41 active sessions, 100% of total activity.

Action

Run SQL Tuning Advisor on the `INSERT` statement with SQL_ID
`"7ucjjhygs7abw"`.

Related Object

SQL statement with SQL_ID `7ucjjhygs7abw`.

Rationale

The SQL spent 100% of its database time on CPU, I/O and Cluster waits.

This part of database time may be improved by the SQL Tuning Advisor.

Rationale

Database time for this SQL was divided as follows: 100% for SQL execution, 0% for parsing, 0% for PL/SQL execution and 0% for Java execution.

Rationale

Top level calls to execute the PL/SQL statement with SQL_ID "ar6t9h4a8sdah" are responsible for 100% of the database time spent on the INSERT statement with SQL_ID "7ucjjhygs7abw".

Related Object

SQL statement with SQL_ID ar6t9h4a8sdah.

Finding 2: Session Connect and Disconnect

Impact is .01 active sessions, 2.6% of total activity.

Session connect and disconnect calls were consuming significant database time.

Recommendation 1: Application Analysis

Estimated benefit is .01 active sessions, 2.6% of total activity.

Action

Investigate application logic for possible reduction of connect and disconnect calls. For example, you might use a connection pool scheme in the middle tier.

Finding 3: Hard Parse

Impact is .01 active sessions, 2.35% of total activity.

Hard parsing of SQL statements was consuming significant database time.

Hard parses due to cursor environment mismatch were not consuming significant database time.

Hard parsing SQL statements that encountered parse errors was not consuming significant database time.

Hard parses due to literal usage and cursor invalidation were not consuming significant database time.

```
The SGA was adequately sized.
No recommendations are available.

~~~~~
~~~~~
~~~~~
Additional Information
-----
Miscellaneous Information
-----
Wait class "Application" was not consuming significant database time.
Wait class "Commit" was not consuming significant database time.
Wait class "Concurrency" was not consuming significant database time.
Wait class "Configuration" was not consuming significant database
time.
CPU was not a bottleneck for the instance.
Wait class "Network" was not consuming significant database time.
Wait class "User I/O" was not consuming significant database time.
1 row selected.
```

The SQL tuning advisor is ran per the recommendations.

```
SQL> VARIABLE t VARCHAR2(20);
BEGIN
:t := DBMS_SQLTUNE.create_tuning_task(sql_id => '&SQL_ID');
DBMS_SQLTUNE.execute_tuning_task(:t);
END;
/SQL> SQL> 2 3 4 5 6 7 8 9
Enter value for sql_id: 8h4k7hum8jgz4
old 3: :t := DBMS_SQLTUNE.create_tuning_task(sql_id => '&SQL_ID');
new 3: :t := DBMS_SQLTUNE.create_tuning_task(sql_id =>
'8h4k7hum8jgz4');
PL/SQL procedure successfully completed.
Elapsed: 00:00:00.67
SQL>
SQL> SET LINES 32767 PAGES 0 TRIMSPOOL ON VERIFY OFF LONG 1000000
LONGC 1000000
SELECT DBMS_SQLTUNE.report_tuning_task (task_name => :t, TYPE =>
'TEXT') FROM DUAL;SQL>
SQL>
```

GENERAL INFORMATION SECTION

```

Tuning Task Name : TASK_394
Tuning Task Owner : SYS
Workload Type : Single SQL Statement
Scope : COMPREHENSIVE
Time Limit(seconds): 1800
Completion Status : COMPLETED
Started at : 10/27/2022 03:09:33
Completed at : 10/27/2022 03:09:34

```

```

Schema Name: SYS
SQL ID : 8h4k7hum8jgz4
SQL Text : SELECT O.ORDER_ID, O.CUSTOMERID, O.ORDER_METHOD,
P.PRODUCT_ID,
P.DESCRIPTION, OT.QUANTITY FROM HANIF.PRODUCT P,
HANIF.ORDER_ITEM OT, HANIF.ORDERS O WHERE ROWNUM=1

```

There **are no recommendations to improve the statement.**

As stated by Oracle, the SQL Tuning Advisor does not provide recommendations to improve the query. This is good to report, since the DBA can assume that their query is the most efficient it can be.

Performance Testing Matrix

DB Name	Datetime of Check	No. long running queries	No. long running RMAN Job	No. Blocking Sessions	DB Size, free space, and used space	Free space for tablespace	% multipass execution and PGA status	Shared pool buffer cache status	Outstanding Alerts	Total users and their type	Topmost wait events, DB time, processing time, and wait
DBST670	10/30/2022 5:00 PM	1	1	1	3gb, 1gb free, 2gb used	All tables have space	0%, PGA is ok	Sufficient	No	33 users 30 background, 3 foreground	SQL*Net Message from client DB Time 942 wait time 935, processing time 7

Conclusion

The order data both stored in the virtual machine and imported undergo integration into the Oracle instance implemented on to the Linux machine. Order data is some of the most popular data in the world resulting in order management system popularity. The database administrator is tasked with many steps in order database implementation including planning, importing data, backup, recovery, statistical analysis, optimization, and maintenance. Oracle 19c on Linux is used to host the database instance. Database performance is confirmed and optimized.

References

- Amulya, M. P. (2020). Progression of Order Management System for Digital Marketing. *International Journal of Advanced Research in Computer and Communication Engineering*, 9(4), 109-116. <https://doi.org/10.17148/IJARCCE.2020.9418>
- Oracle. (n.d.a). Oracle CPU used by this session. Oracle Consulting, Oracle Support and Oracle Training by BC Oracle Consulting.
https://www.dba-oracle.com/m_cpu_used_by_this_session.htm
- Burleson. (n.d.b). Oracle redo wastage. Oracle Consulting, Oracle Support and Oracle Training by BC Oracle Consulting. https://www.dba-oracle.com/m_redo_wastage.htm

Burleson, D. (n.d.c). Oracle sorts memory. Oracle Consulting, Oracle Support and Oracle Training by BC Oracle Consulting. https://www.dba-oracle.com/m_sorts_memory.htm

Mockaroo. (2022). *Random Data Generator and API Mocking Tool*. <https://www.mockaroo.com/>

Oracle. (2019, July 31). Automatic database diagnostic monitor (ADDM) in Oracle database 10g. oracle-base.com.

<https://oracle-base.com/articles/10g/automatic-database-diagnostic-monitor-10g>

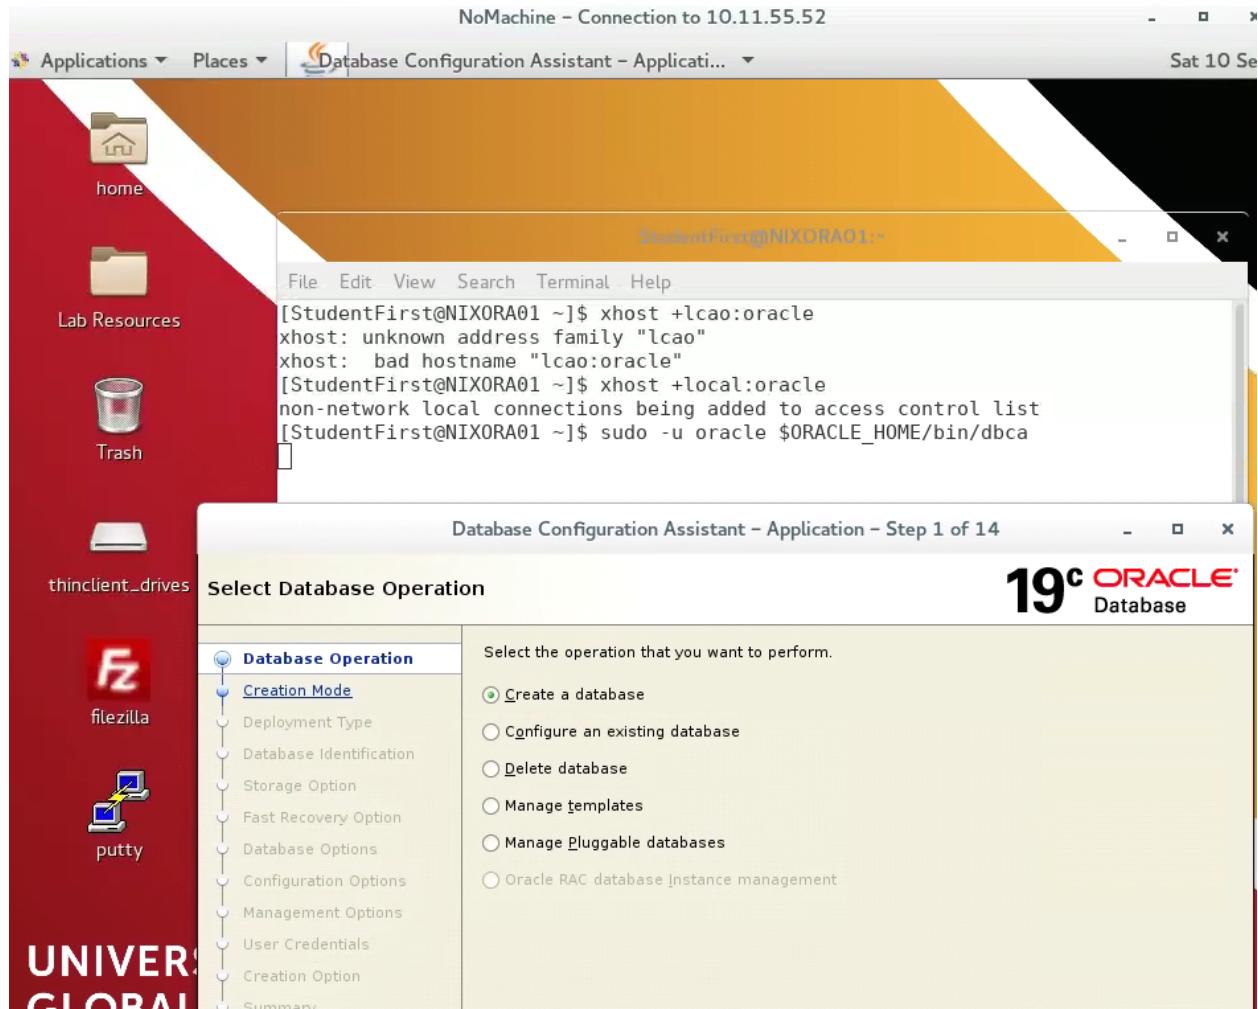
Oracle. (n.d.). V\$sysstat.

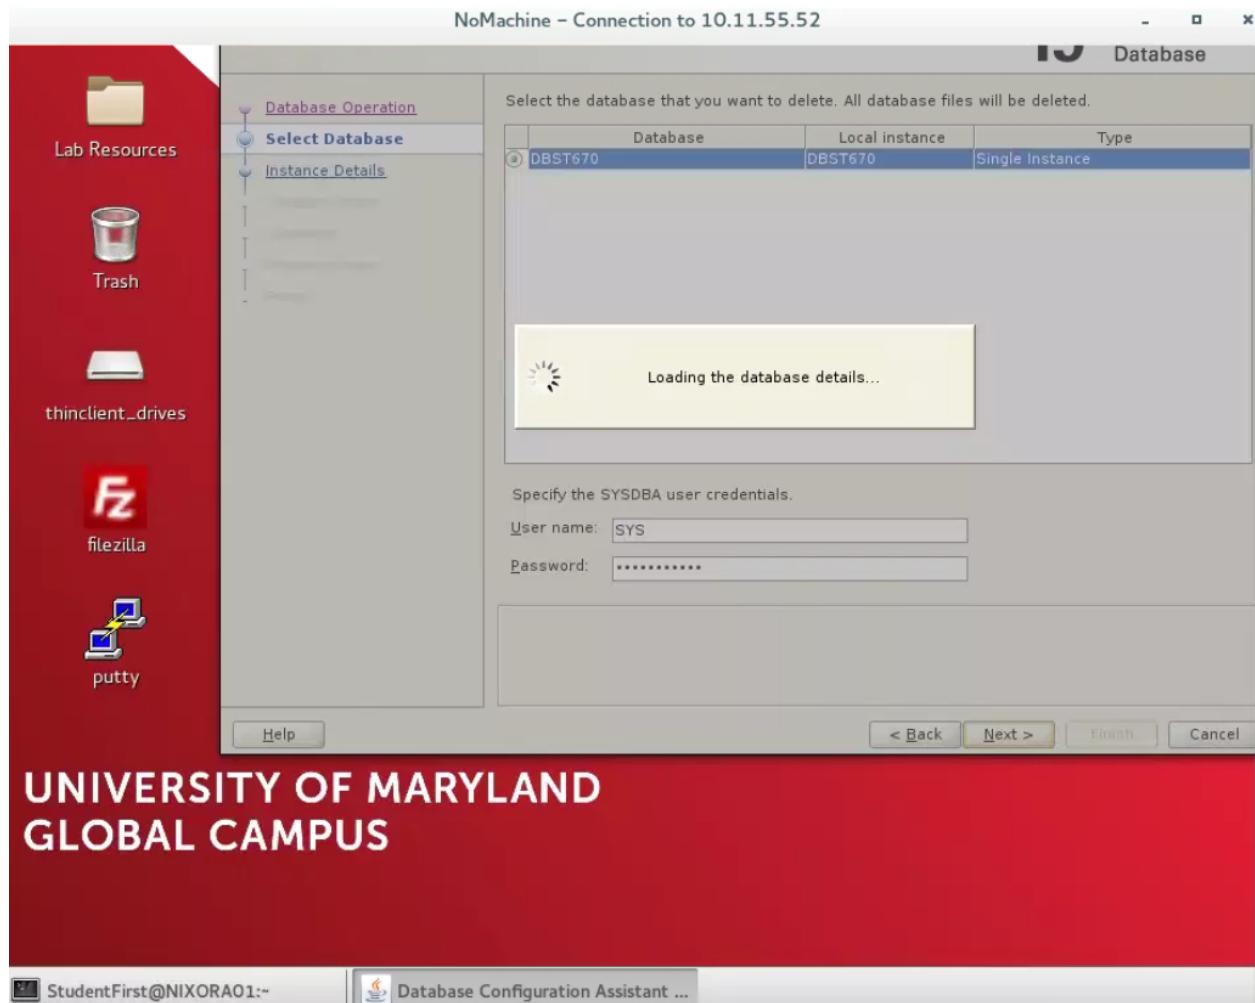
https://docs.oracle.com/cd/E18283_01/server.112/e17110/dynviews_3093.htm

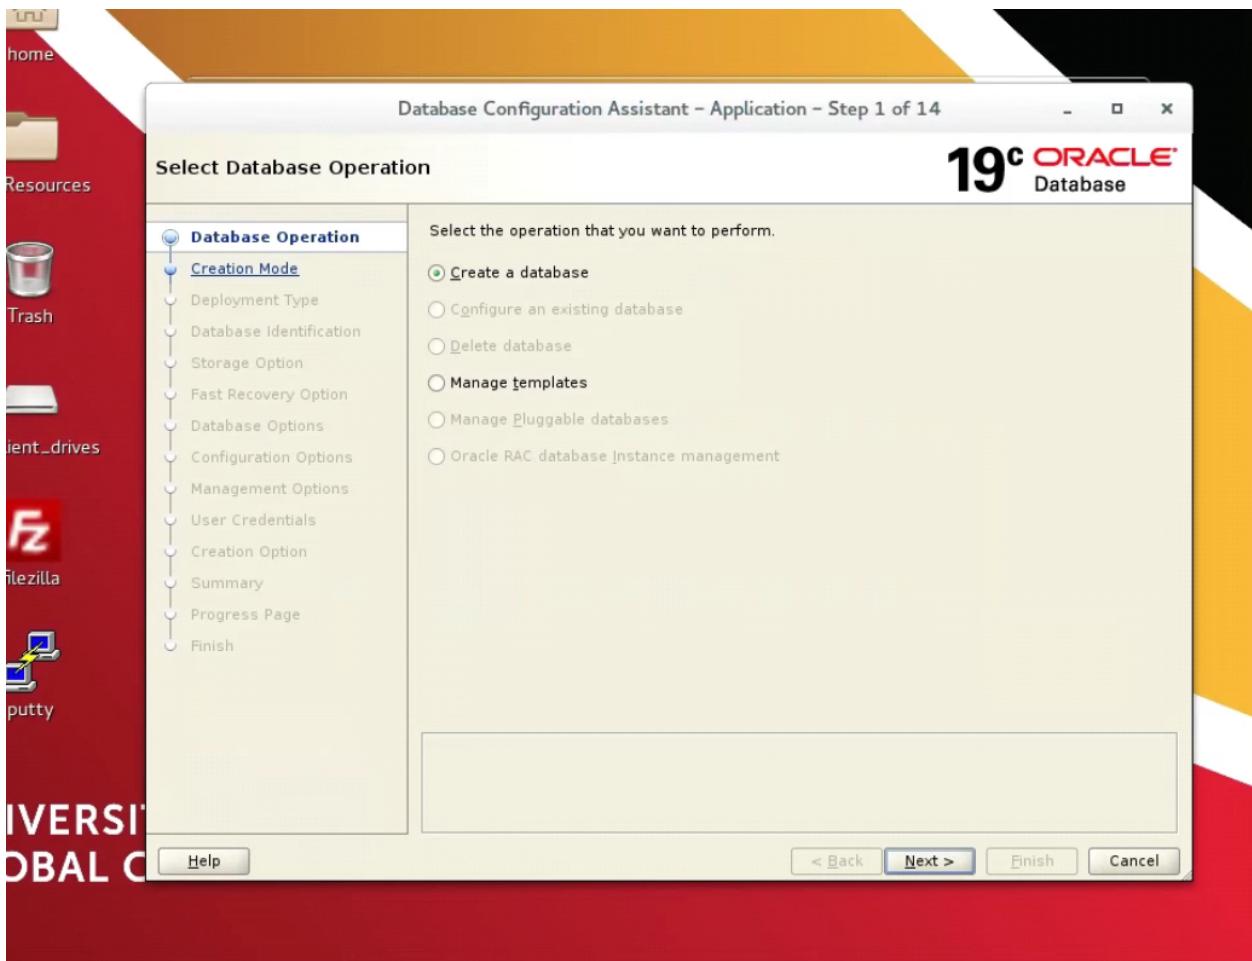
Index

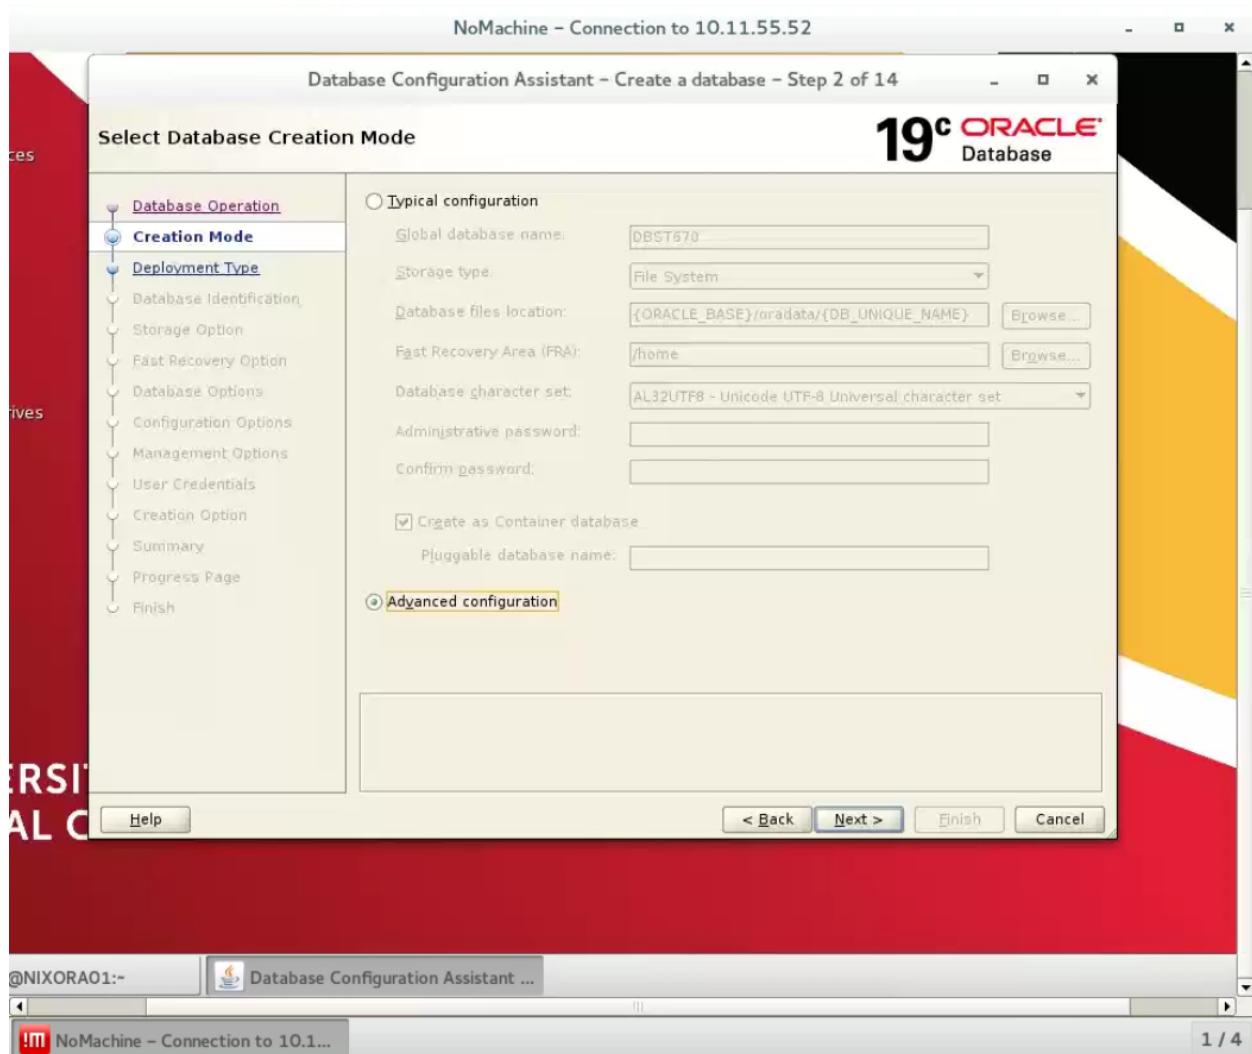
The following is the DDL script generated using Oracle Data Modeler:

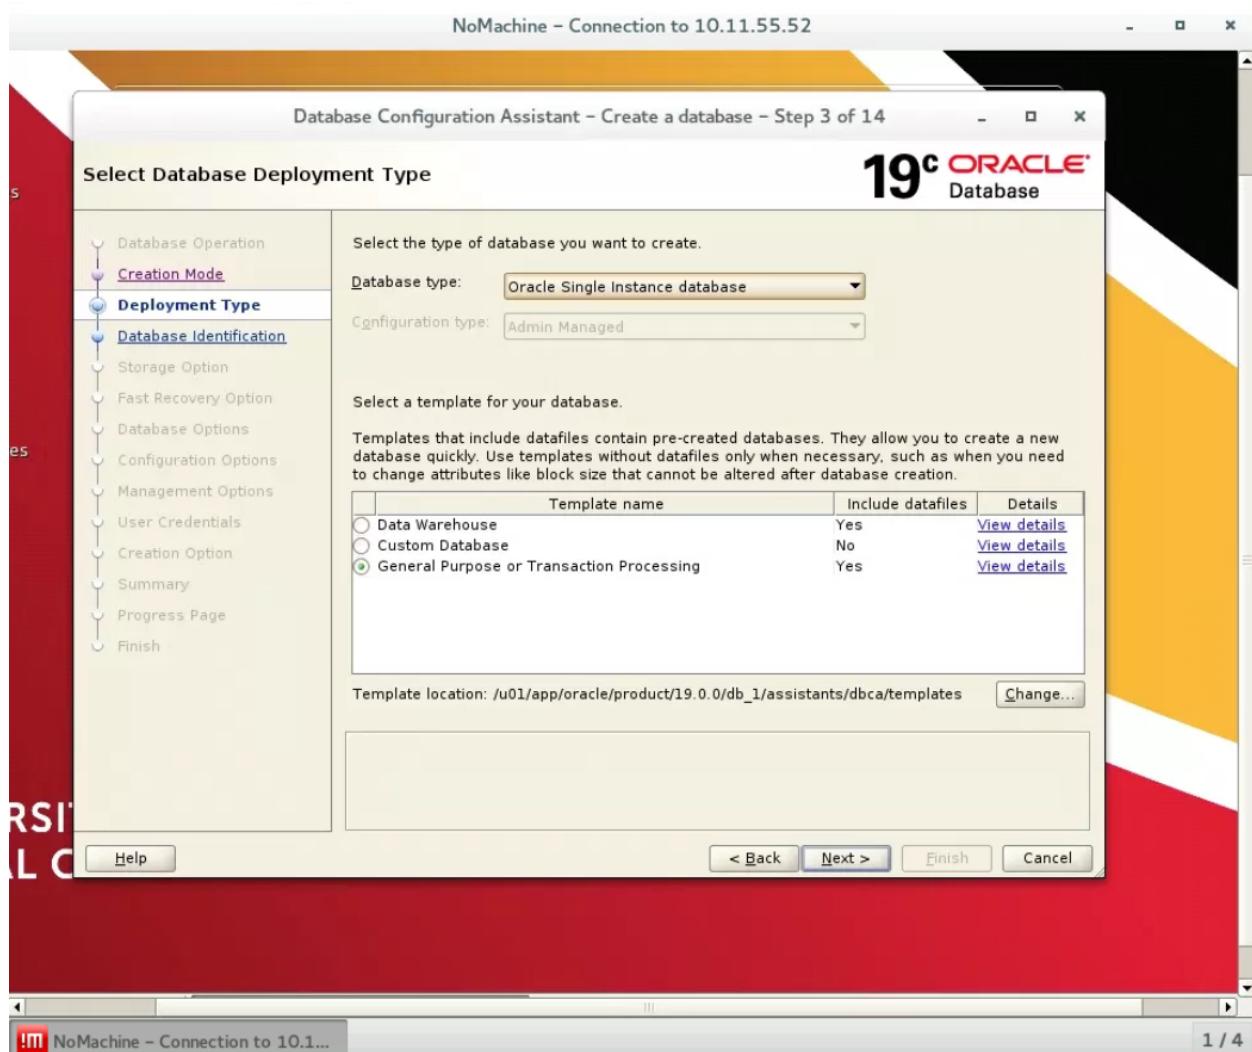
Database Configuration Assistant Wizard Steps

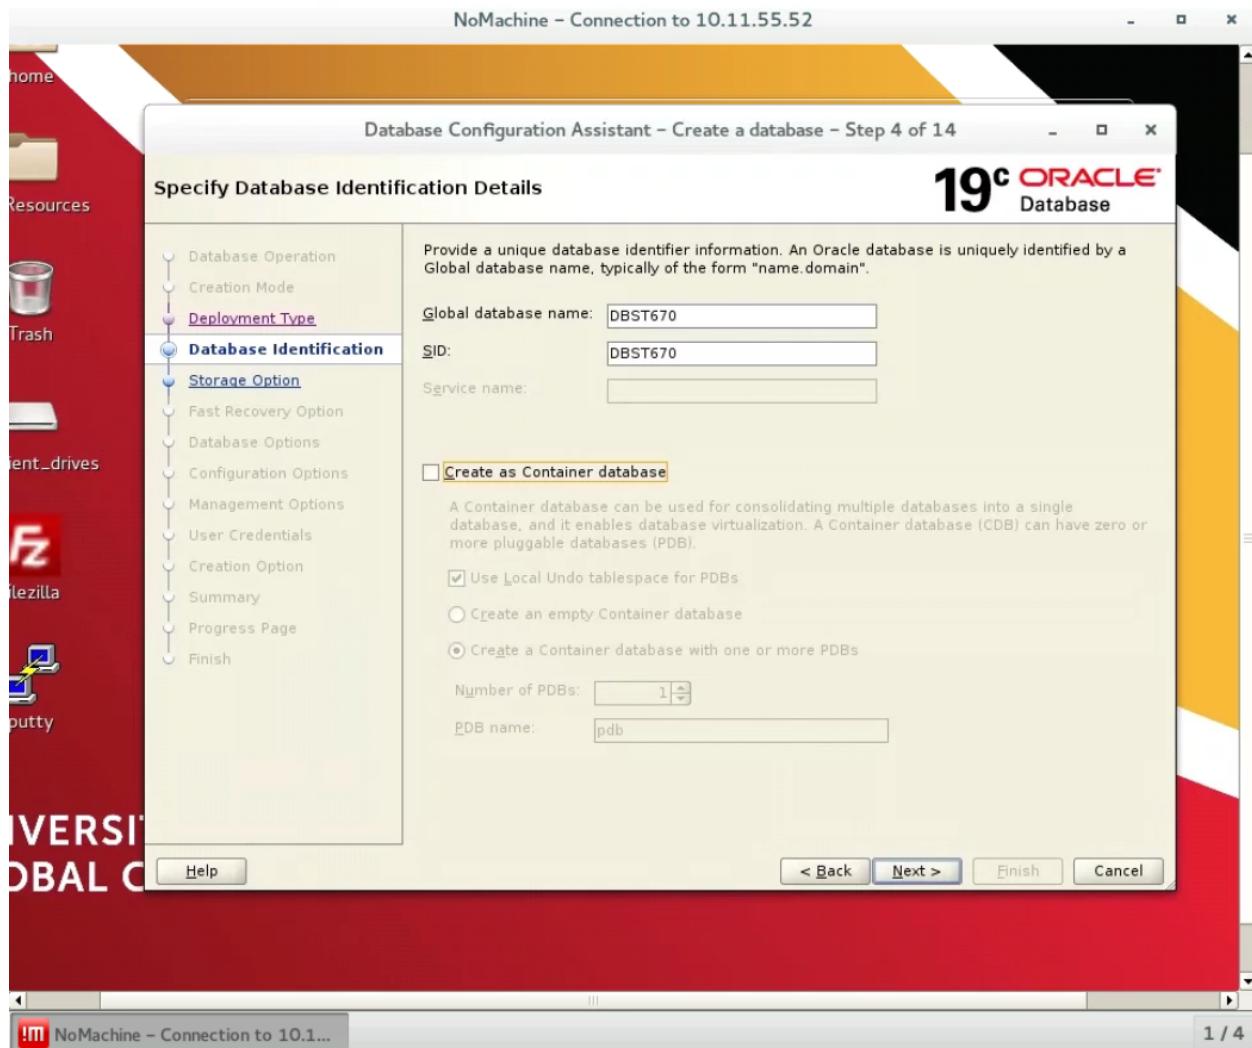


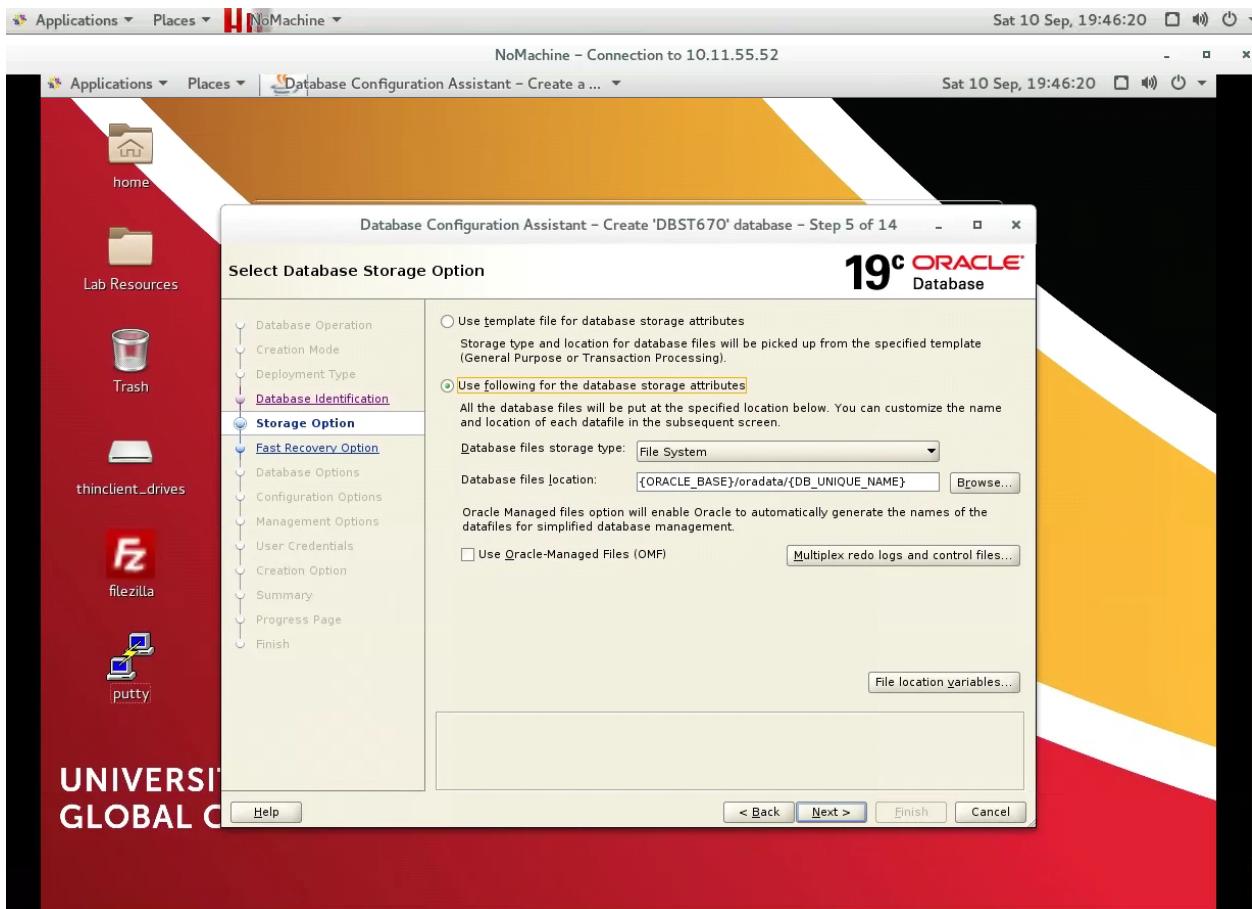


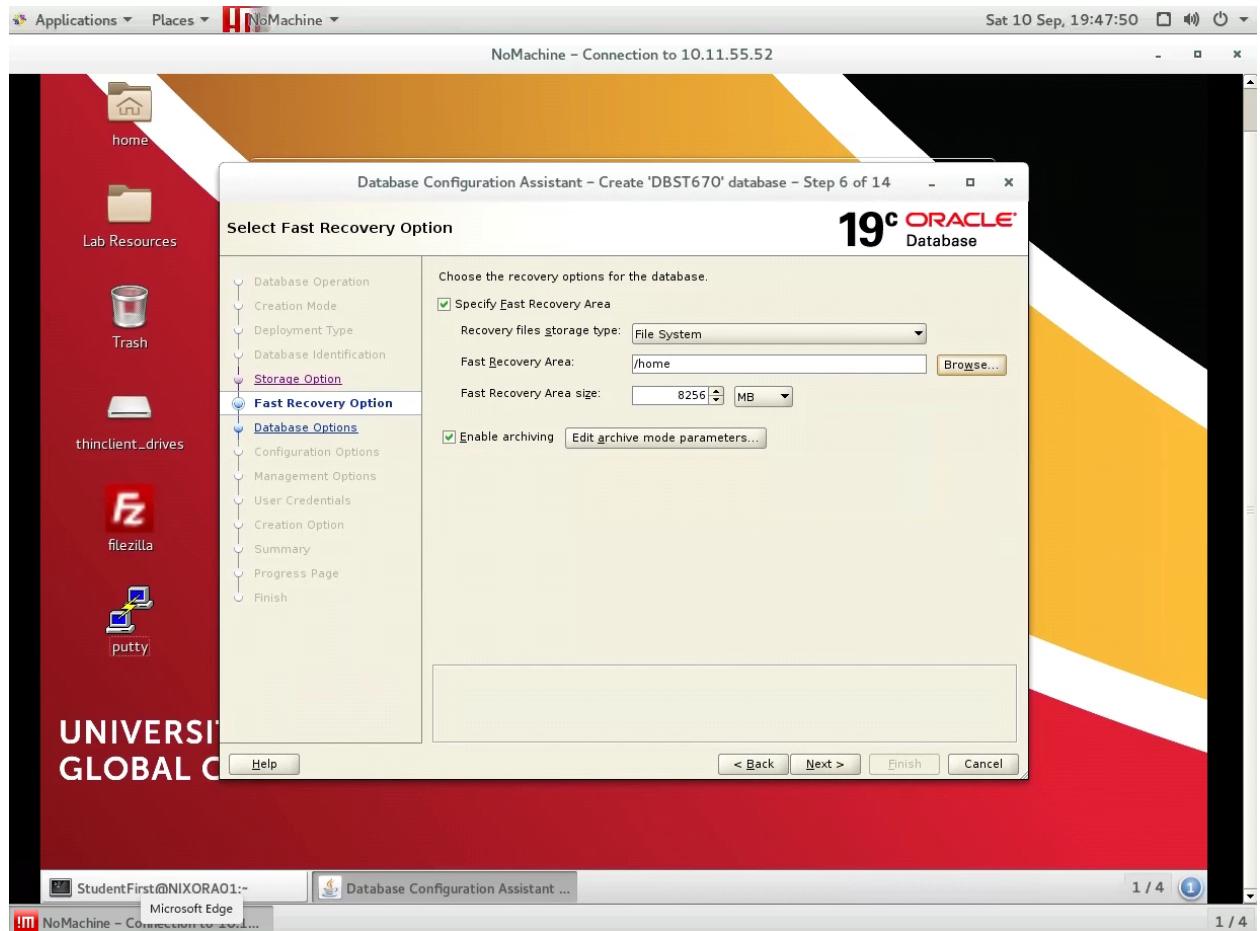


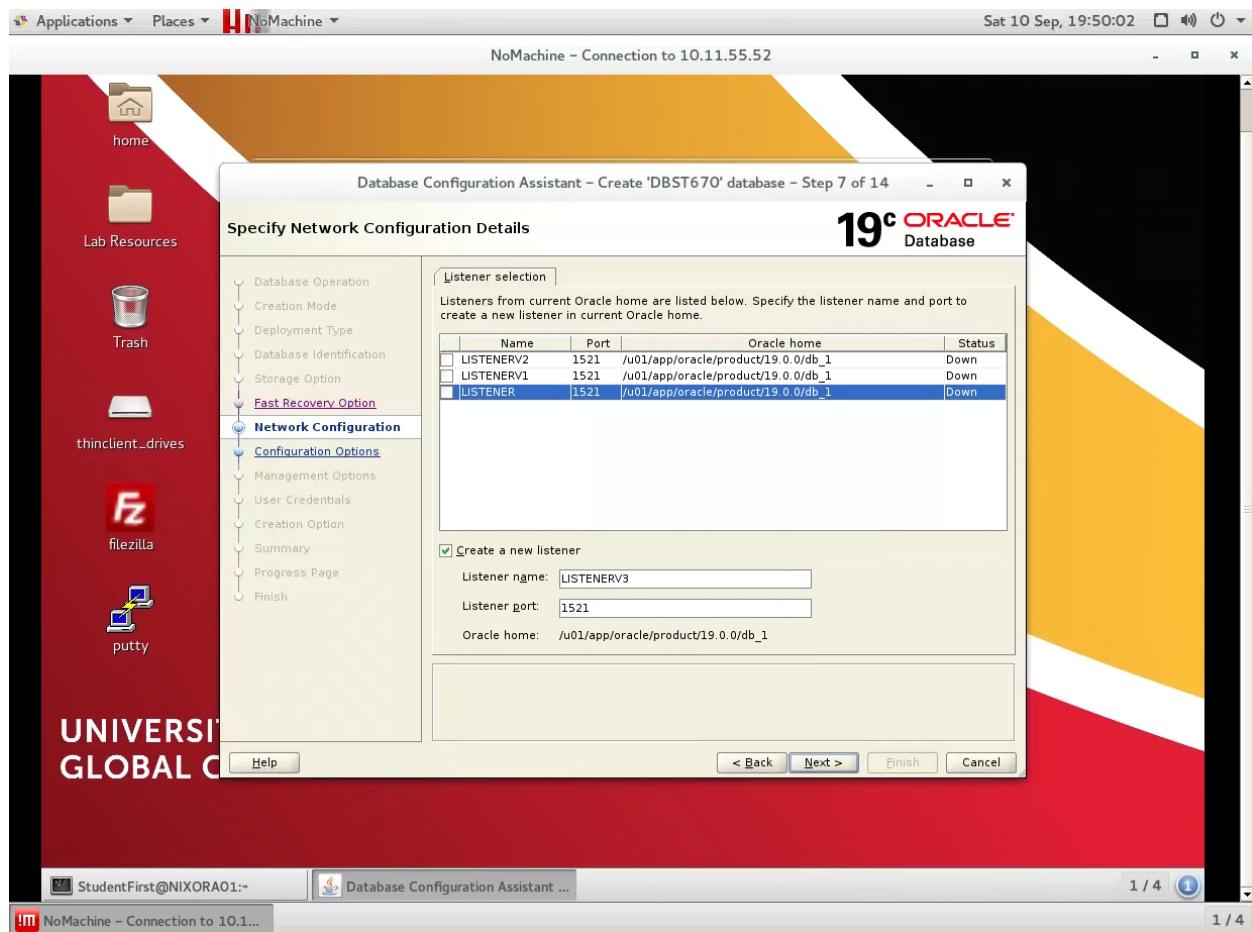


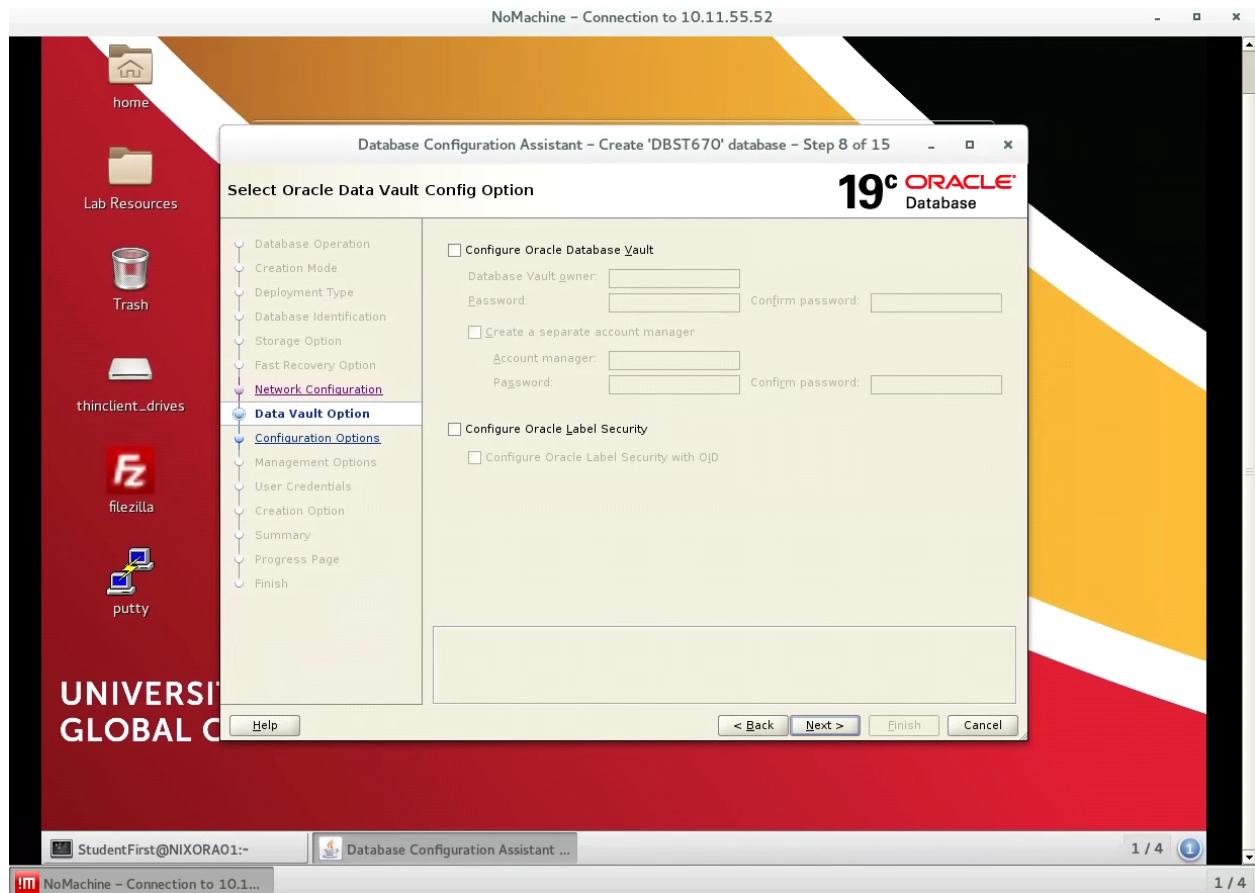


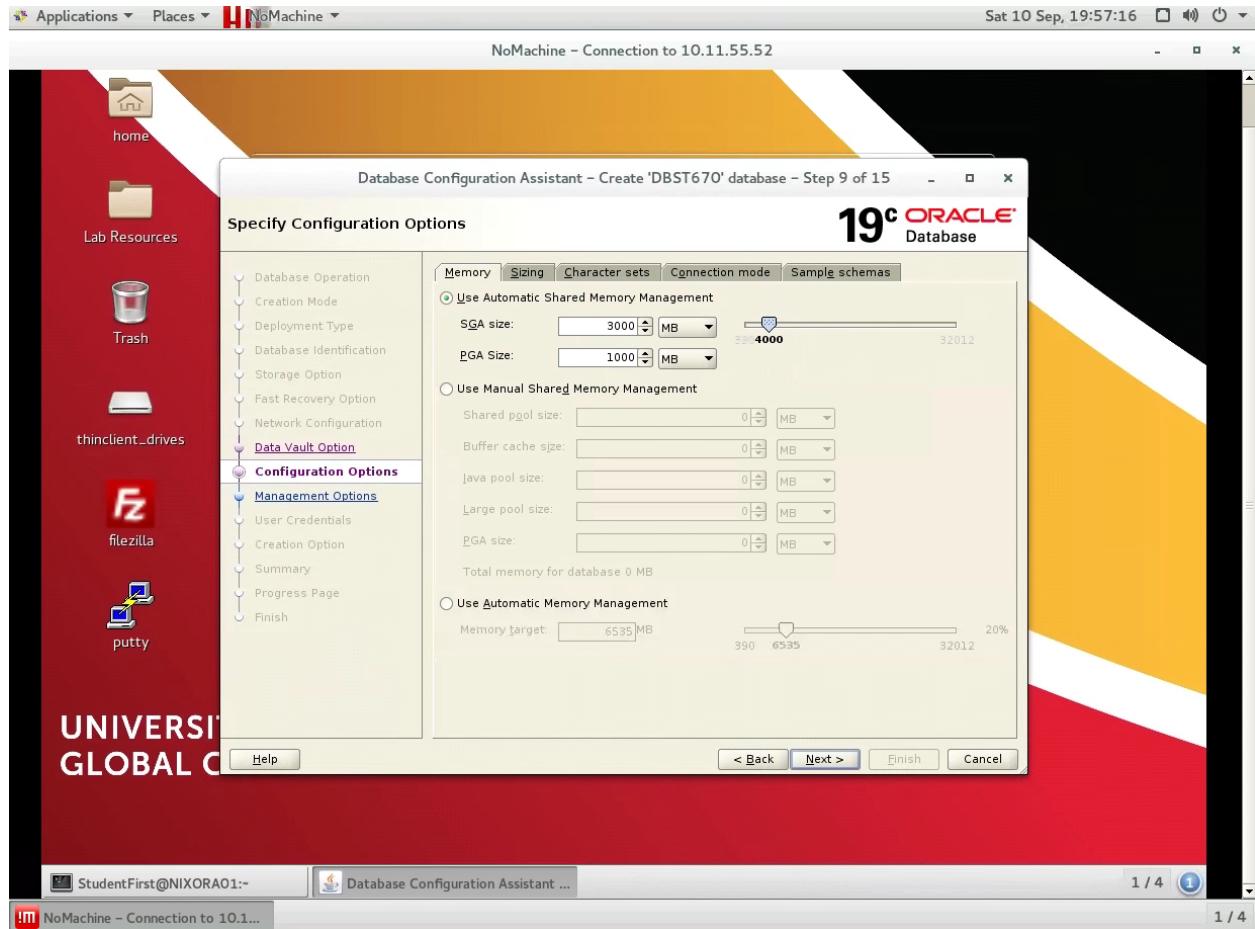


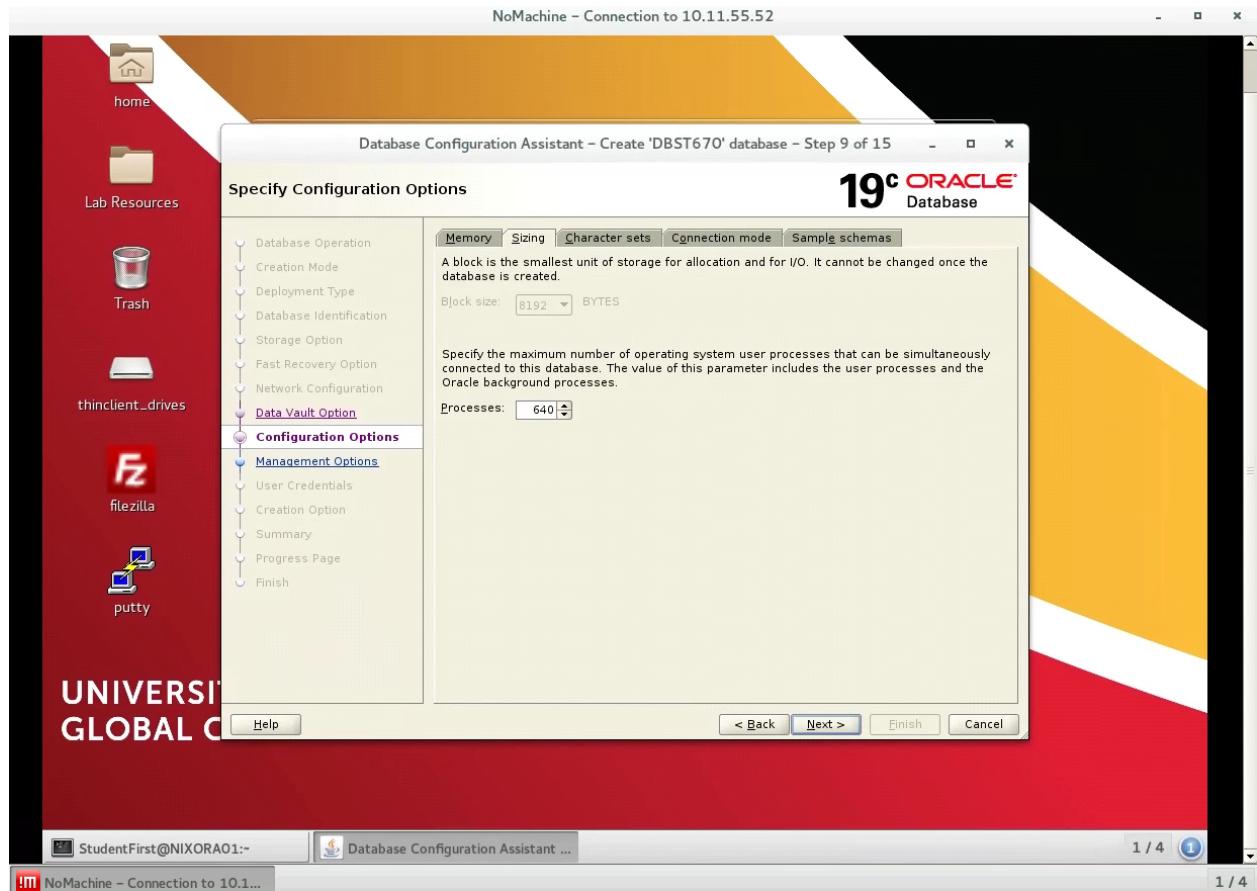


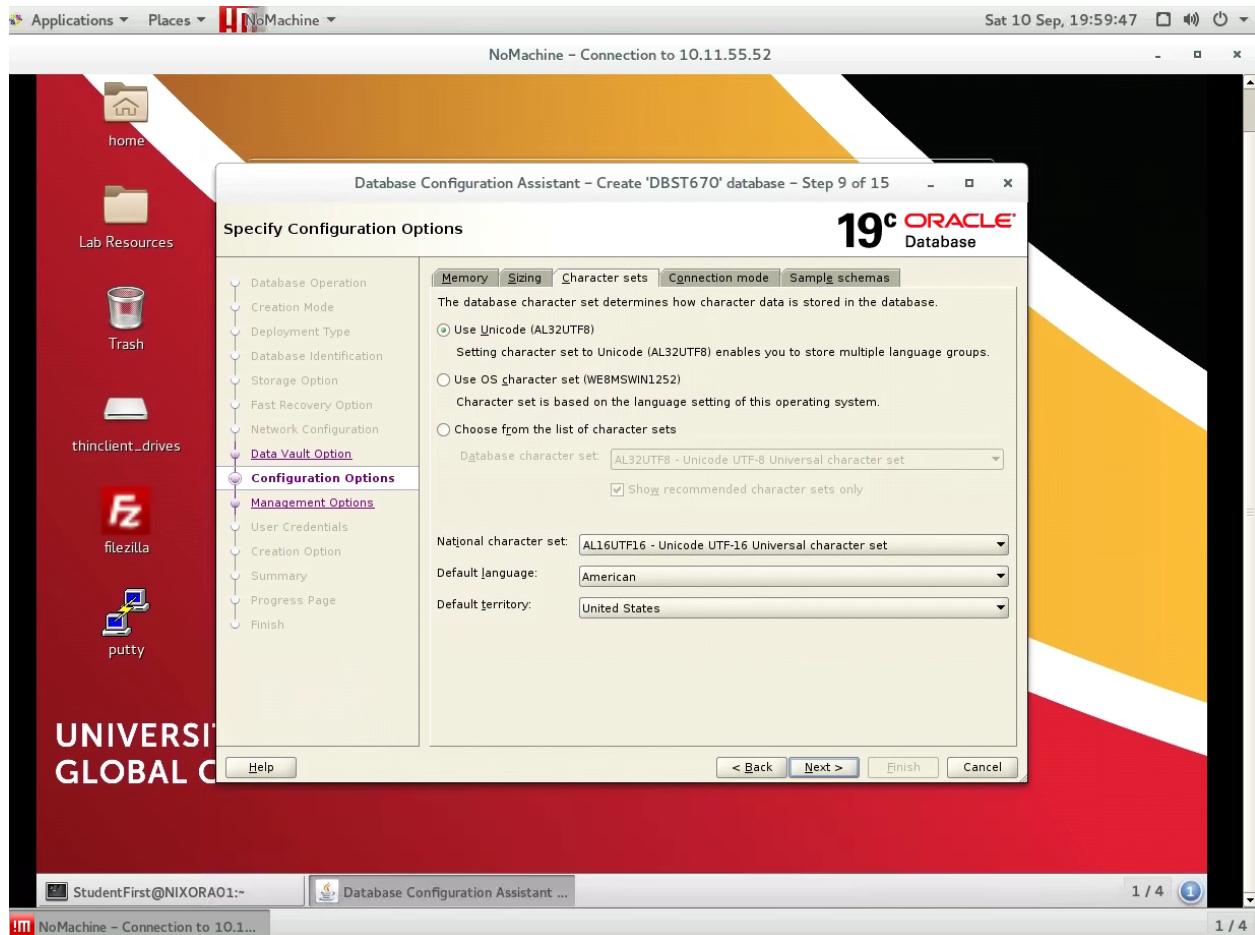


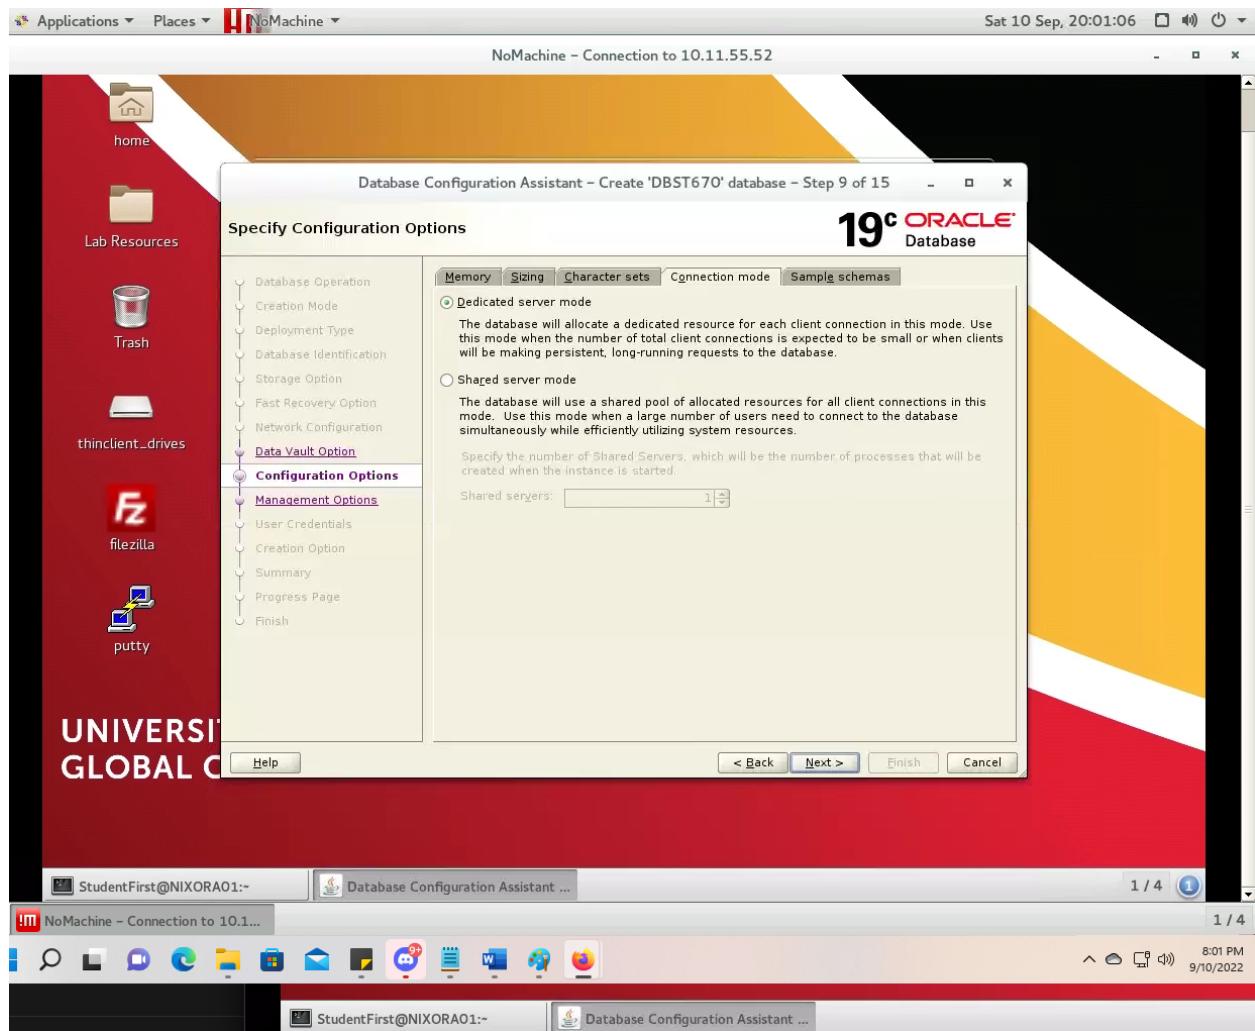


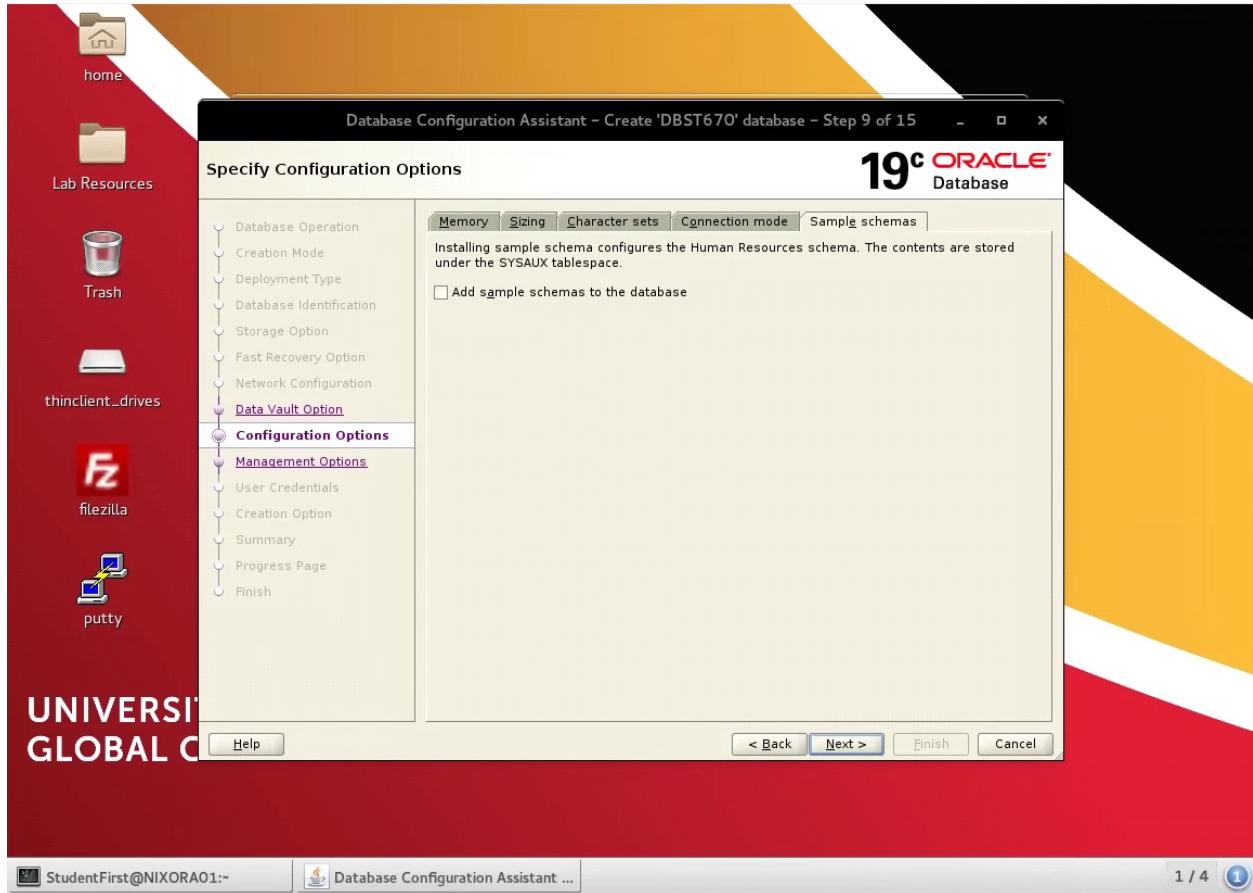


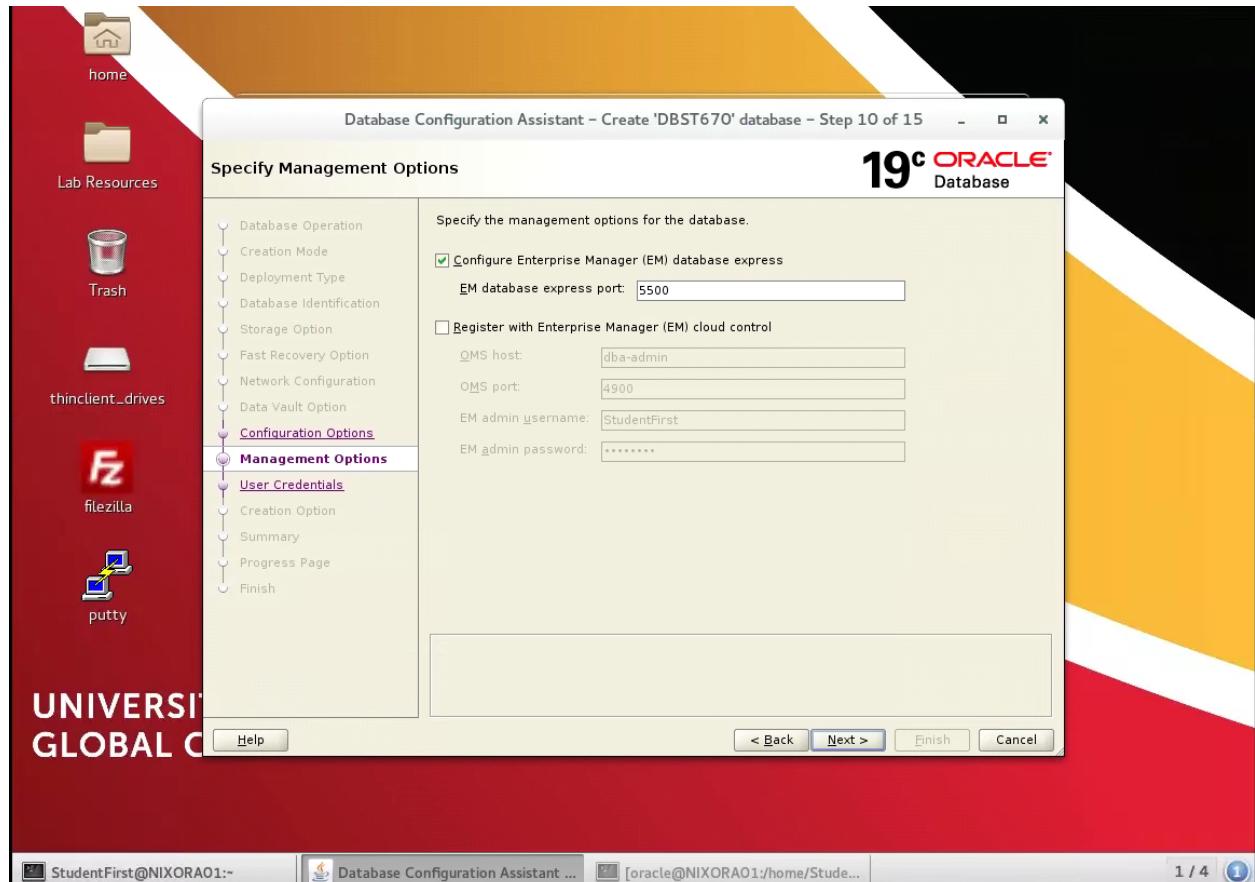


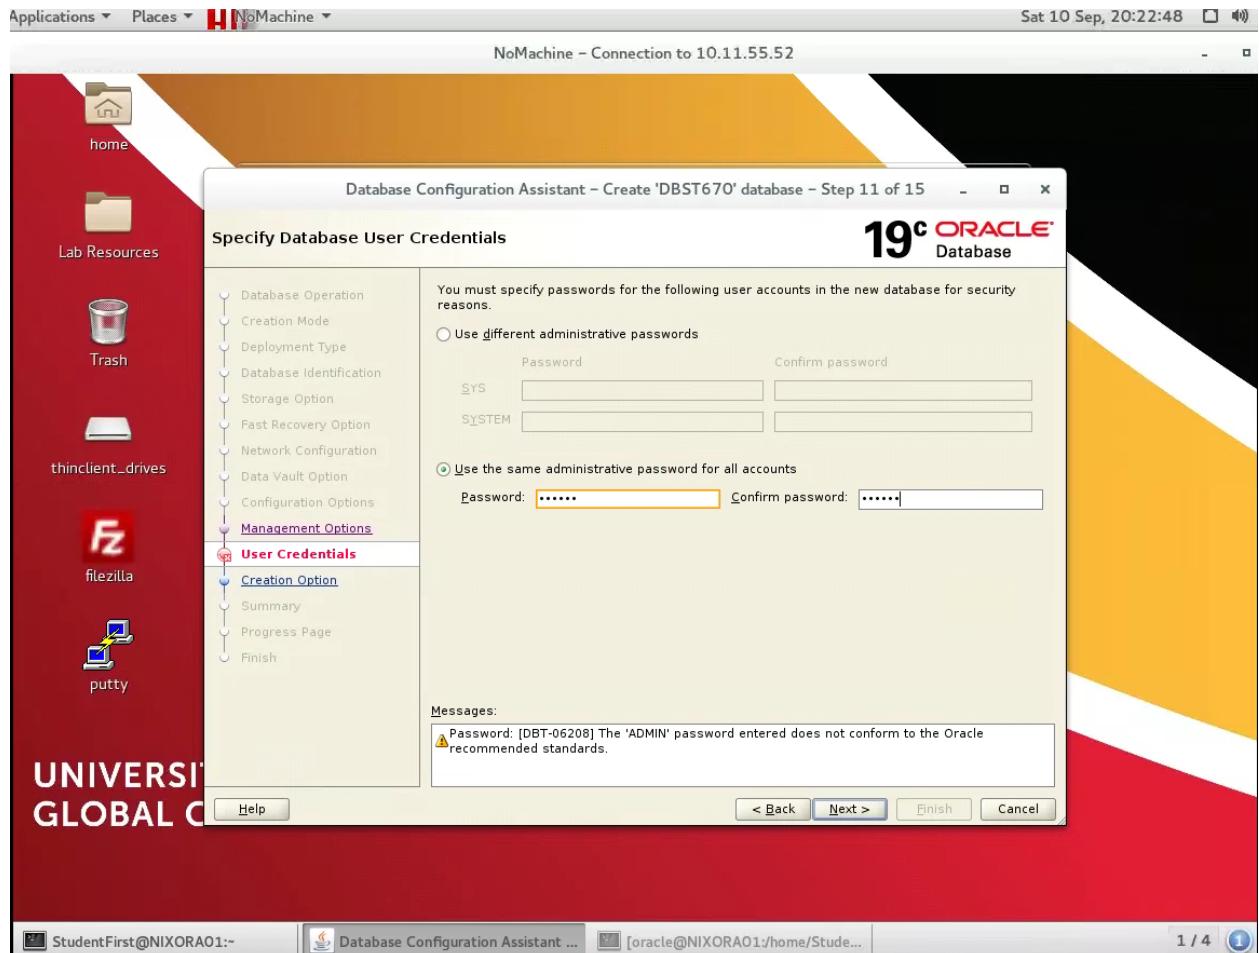


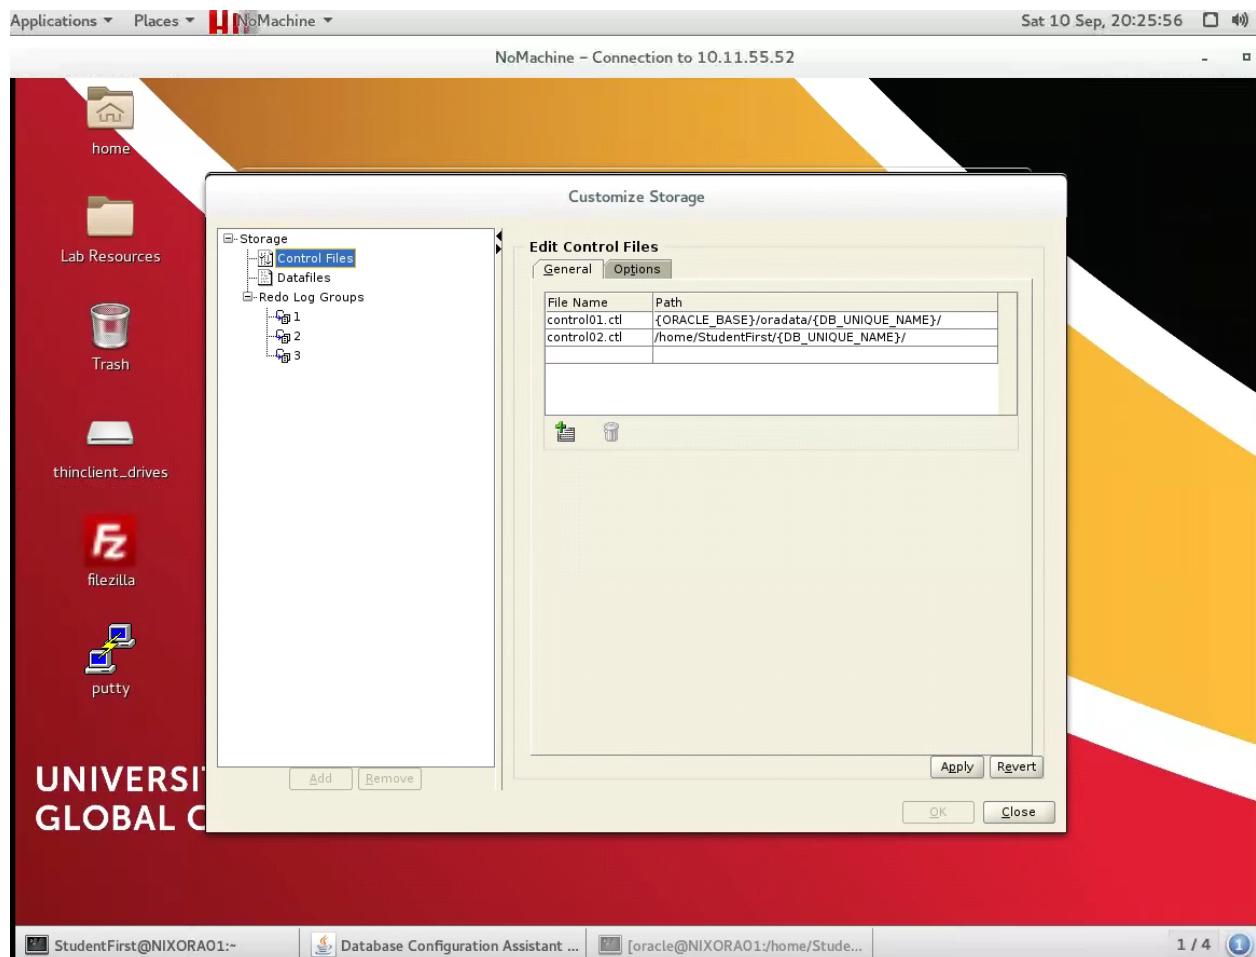


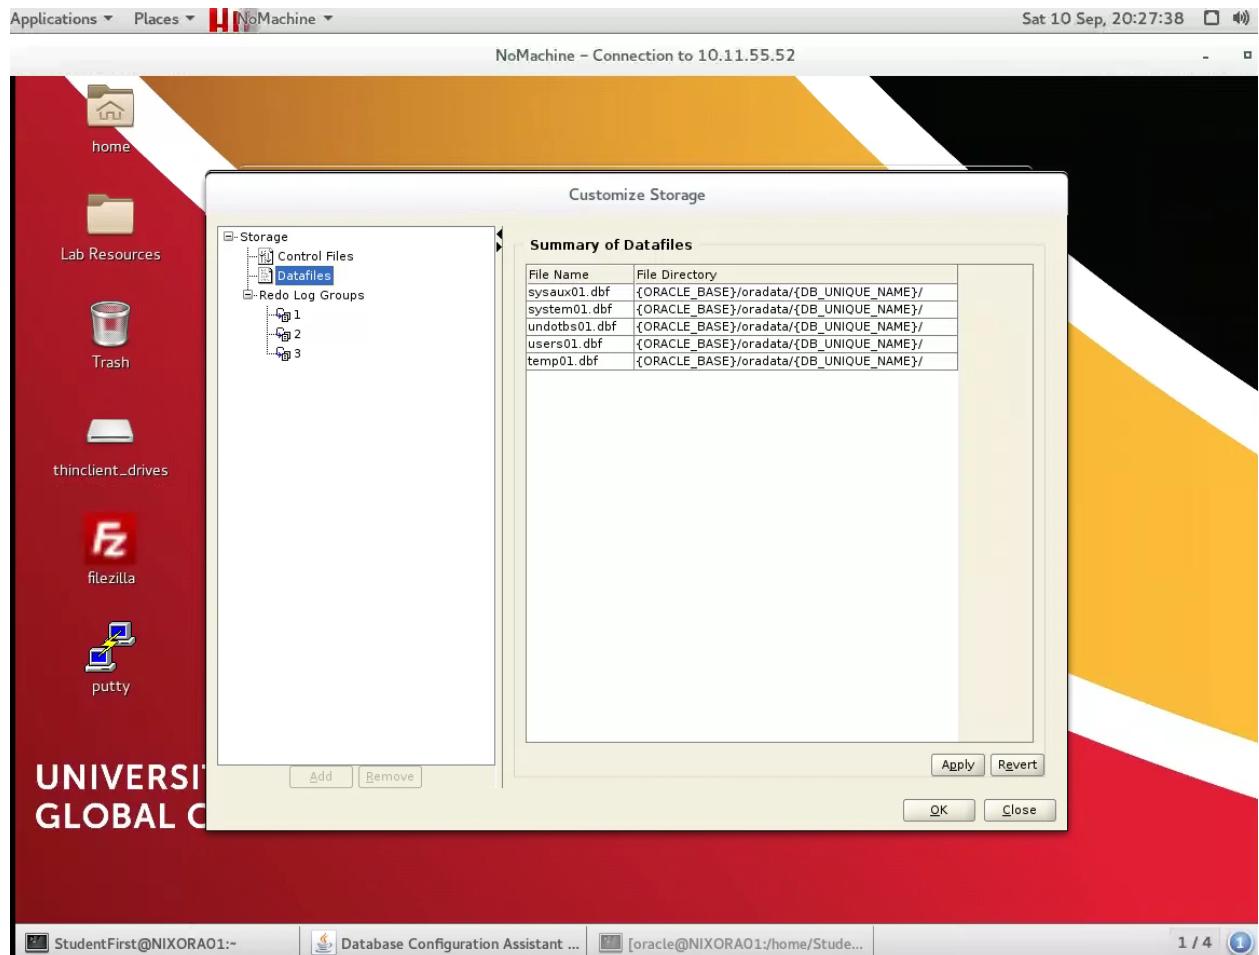


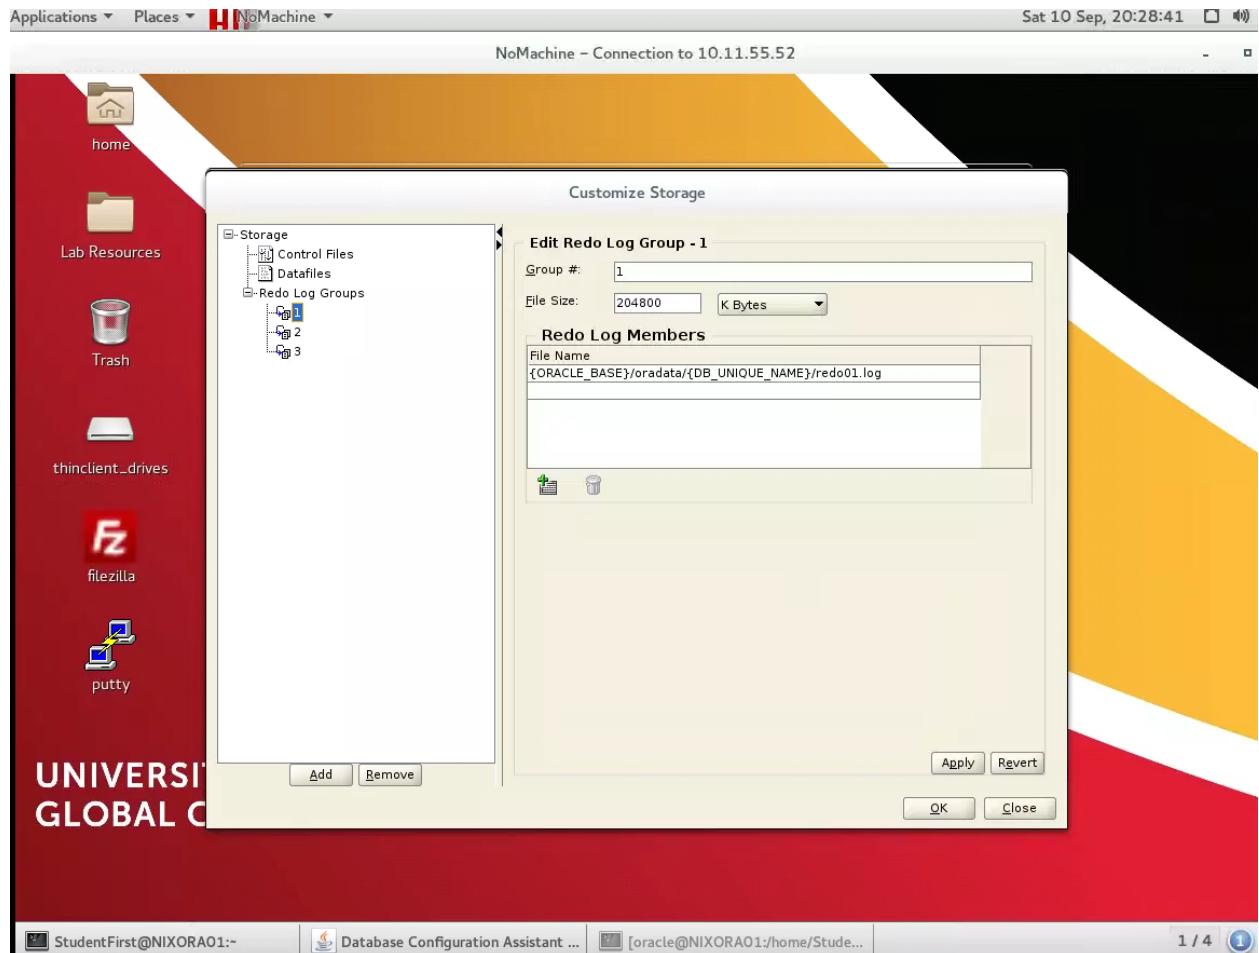


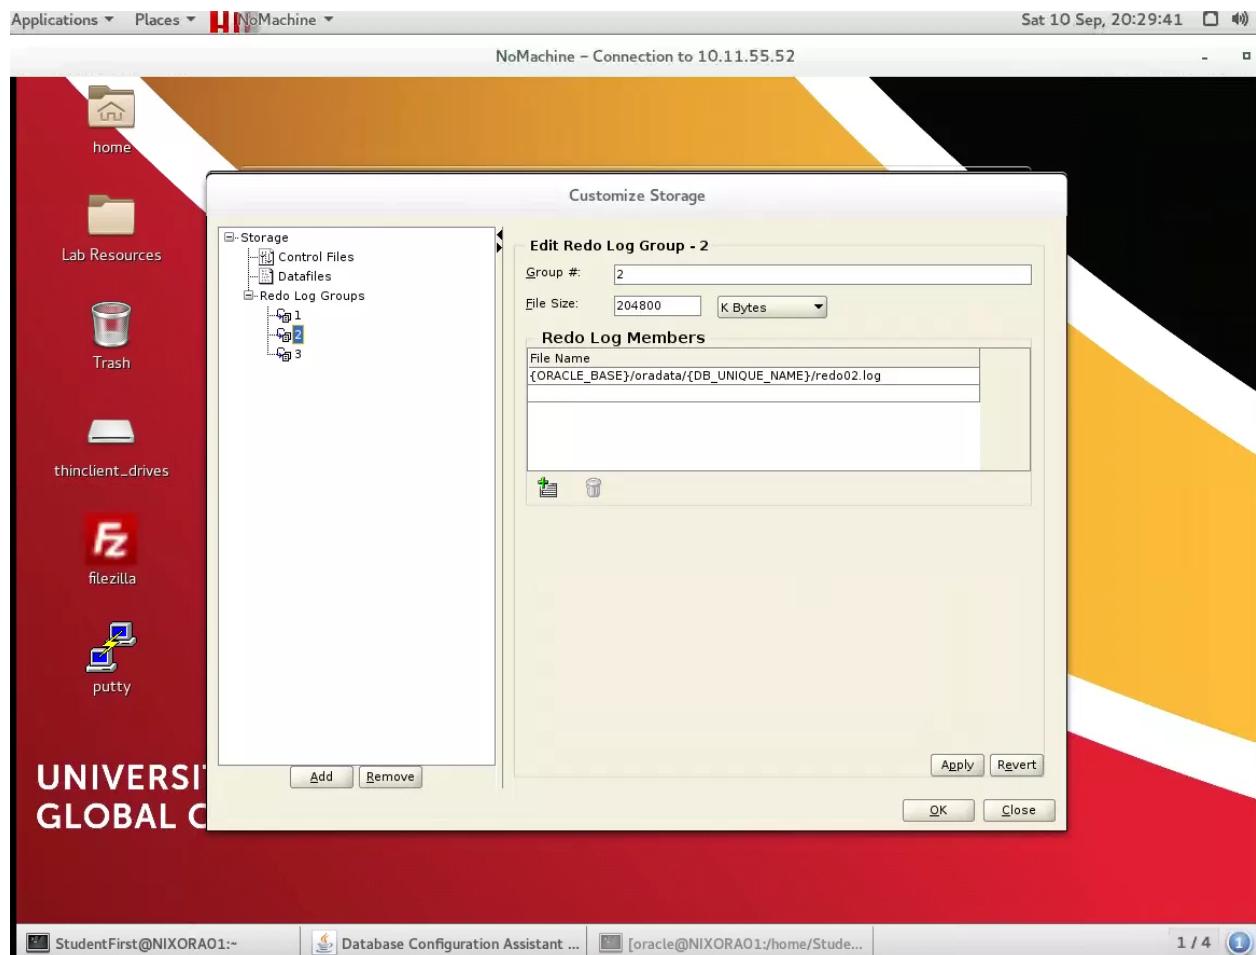


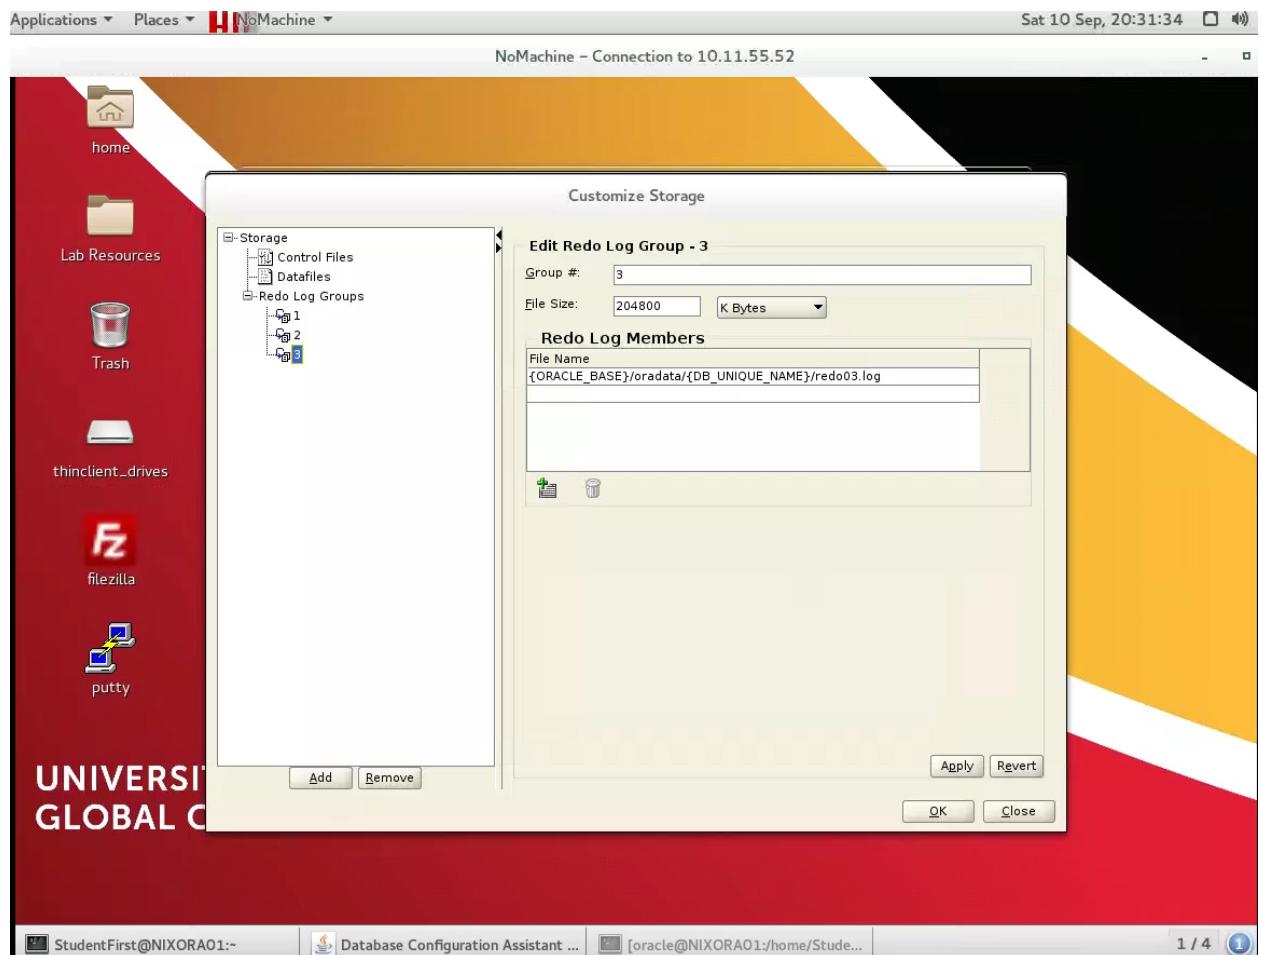


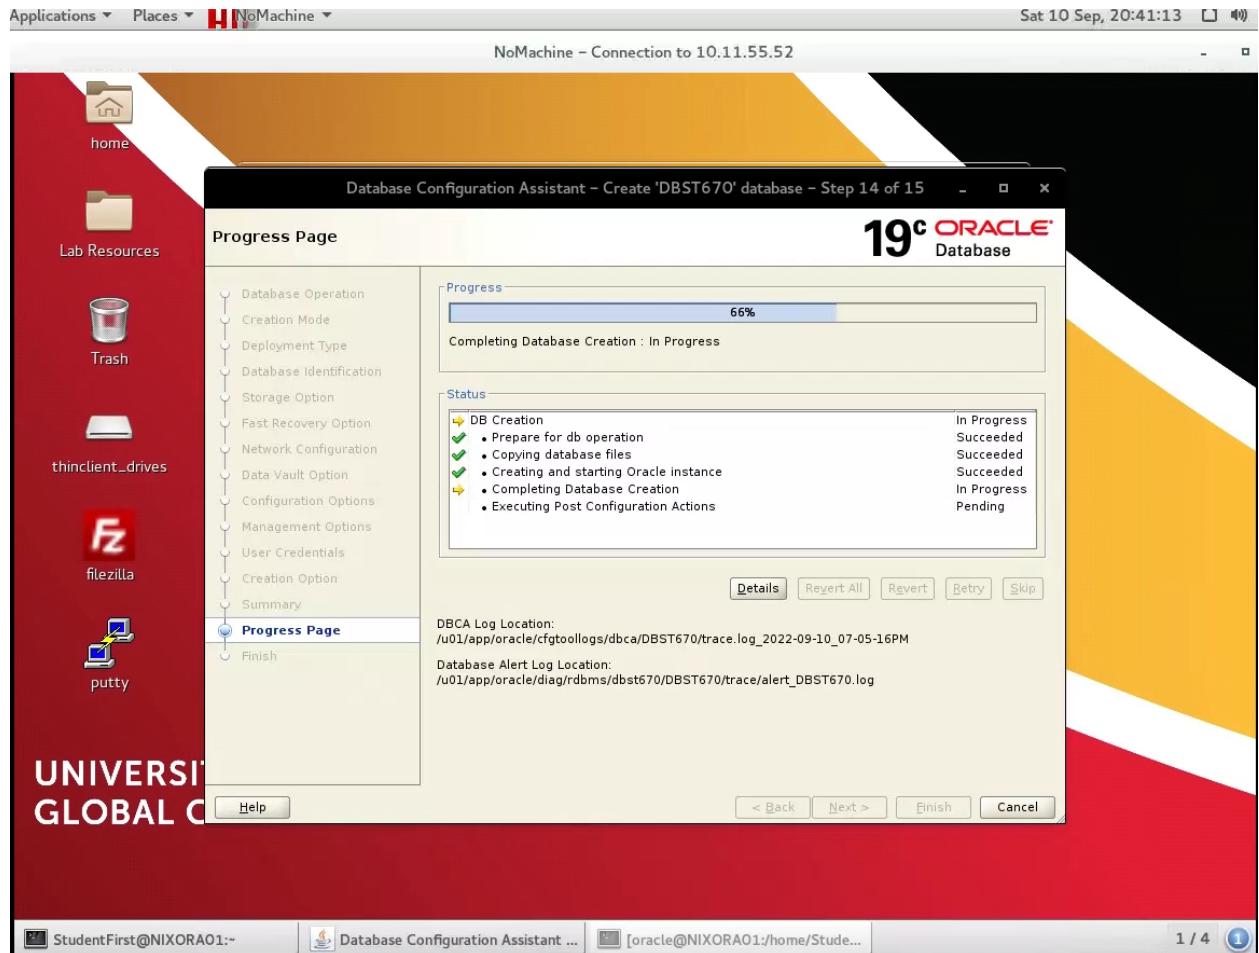


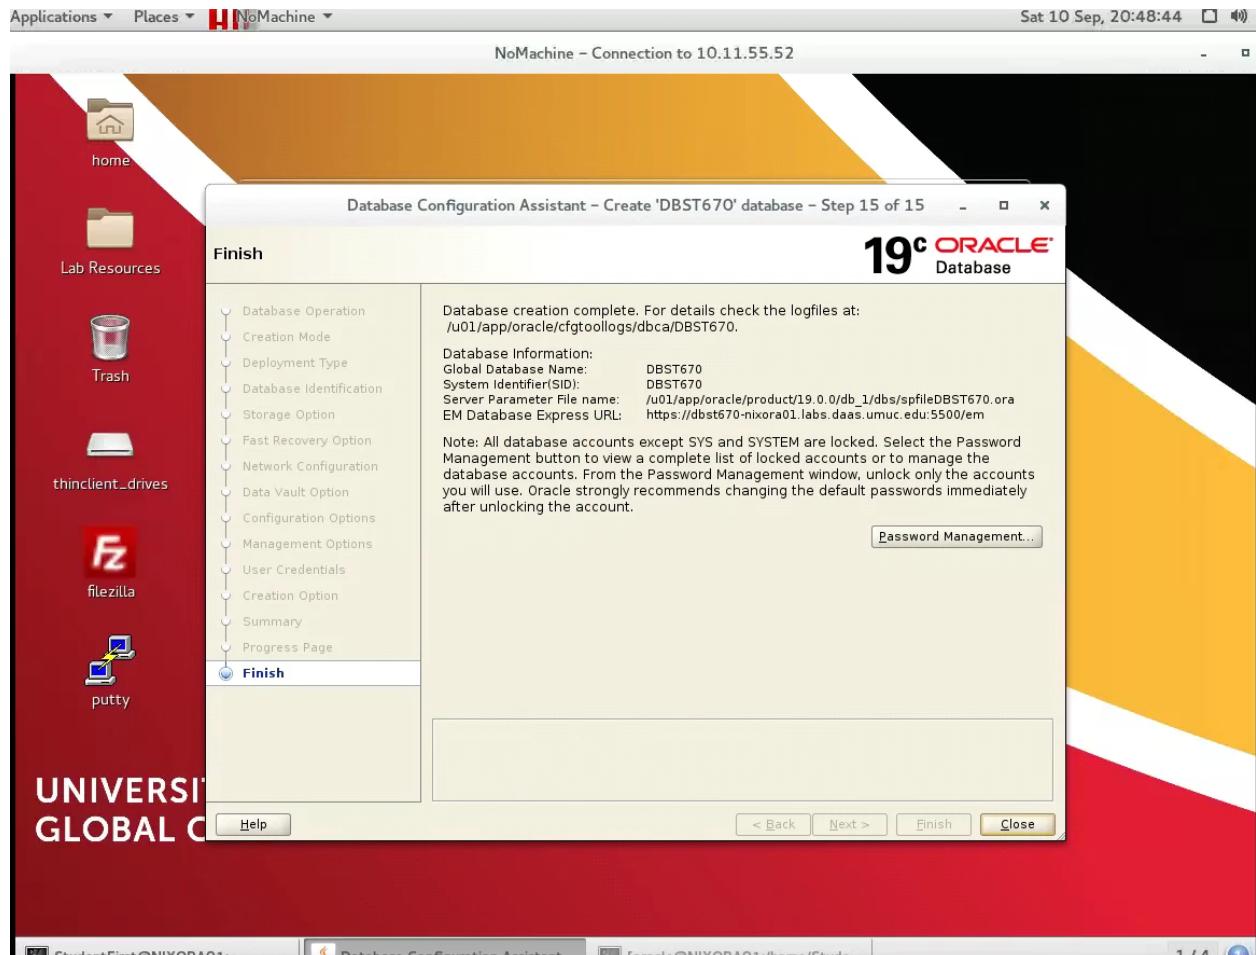


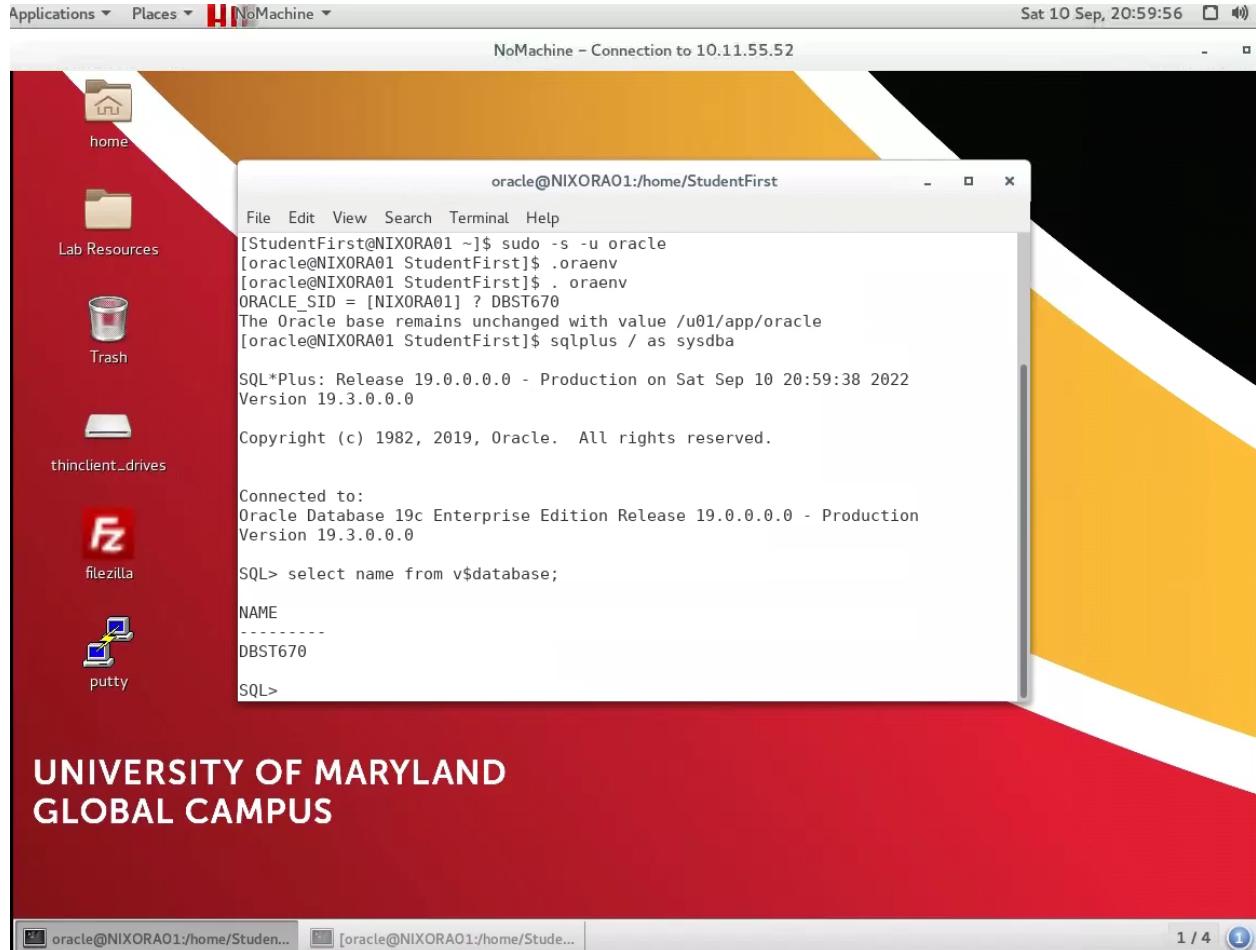












UNIVERSITY OF MARYLAND
GLOBAL CAMPUS

The screenshot shows a terminal window titled "NoMachine – Connection to 10.11.55.52" running on a Linux system. The window title bar includes the connection details and the current date and time: "Sat 10 Sep, 21:14:32". The terminal window has a menu bar with "File", "Edit", "View", "Search", "Terminal", and "Help". The main pane displays the following RMAN command output:

```
Recovery Manager complete.
[oracle@NIXORA01 StudentFirst]$ rman target /
Recovery Manager: Release 19.0.0.0.0 - Production on Sat Sep 10 21:14:00 2022
Version 19.3.0.0.0

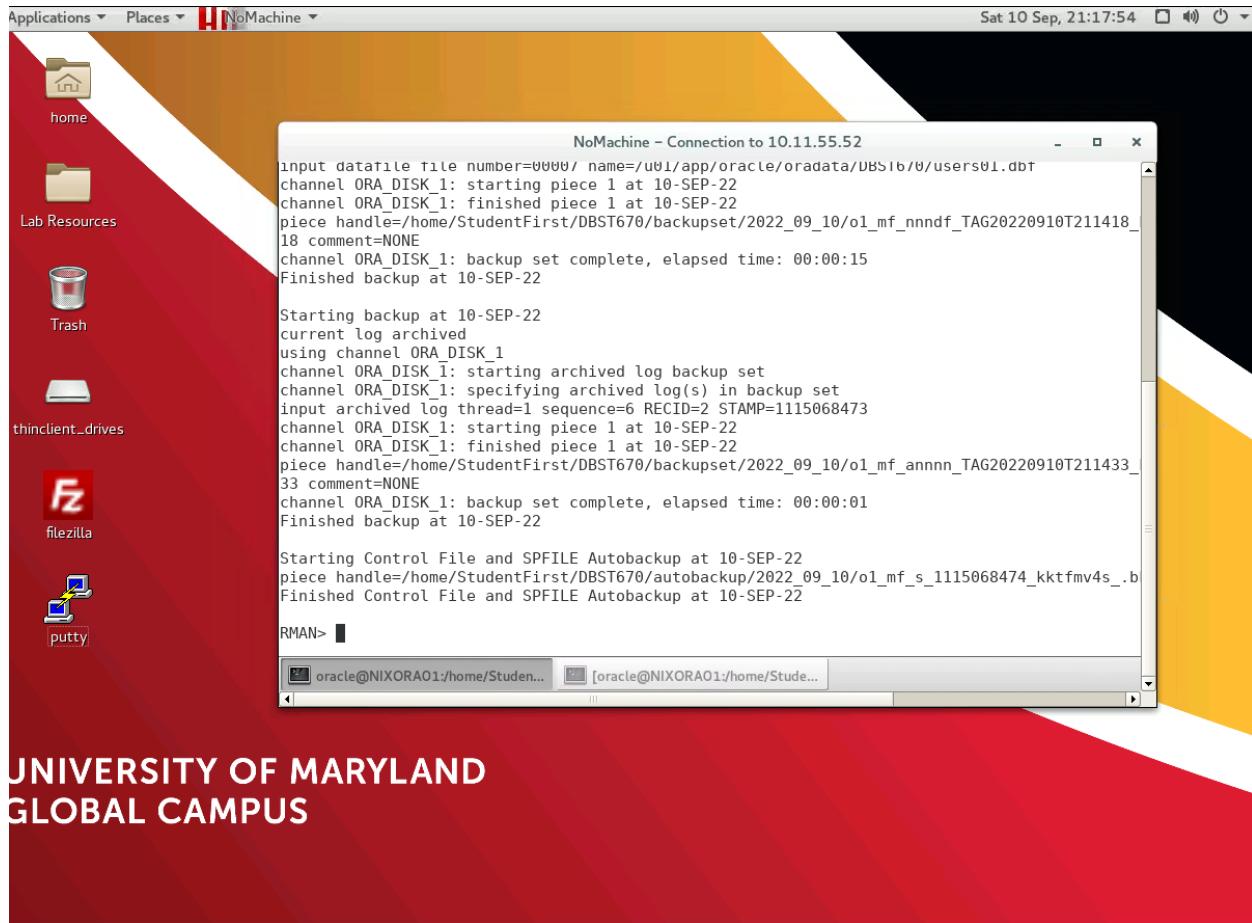
Copyright (c) 1982, 2019, Oracle and/or its affiliates. All rights reserved.

connected to target database: DBST670 (DBID=990969903)

RMAN> backup database plus archivelog;

Starting backup at 10-SEP-22
current log archived
using target database control file instead of recovery catalog
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=382 device type=DISK
channel ORA_DISK_1: starting archived log backup set
channel ORA_DISK_1: specifying archived log(s) in backup set
input archived log thread=1 sequence=5 RECID=1 STAMP=1115068454
channel ORA_DISK_1: starting piece 1 at 10-SEP-22
channel ORA_DISK_1: finished piece 1 at 10-SEP-22
piece handle=/home/StudentFirst/DBST670/backupset/2022_09_10/o1_mf_annnn_TAG20220910T211414_kktfm71z_.bkp tag=TAG20220910T211414 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:03
Finished backup at 10-SEP-22

Starting backup at 10-SEP-22
using channel ORA_DISK_1
channel ORA_DISK_1: starting full datafile backup set
channel ORA_DISK_1: specifying datafile(s) in backup set
input datafile file number=00001 name=/u01/app/oracle/oradata/DBST670/system01.dbf
input datafile file number=00003 name=/u01/app/oracle/oradata/DBST670/sysaux01.dbf
input datafile file number=00004 name=/u01/app/oracle/oradata/DBST670/undotbs01.dbf
input datafile file number=00007 name=/u01/app/oracle/oradata/DBST670/users01.dbf
channel ORA_DISK_1: starting piece 1 at 10-SEP-22
```



UNIVERSITY OF MARYLAND
GLOBAL CAMPUS

```
NoMachine - Connection to 10.11.55.52

Starting backup at 10-SEP-22
current log archived
using channel ORA_DISK_1
channel ORA_DISK_1: starting archived log backup set
channel ORA_DISK_1: specifying archived log(s) in backup set
input archived log thread=1 sequence=6 RECID=2 STAMP=1115068473
channel ORA_DISK_1: starting piece 1 at 10-SEP-22
channel ORA_DISK_1: finished piece 1 at 10-SEP-22
piece handle=/home/StudentFirst/DBST670/backupset/2022_09_10/o1_mf_annnn_TAG20220910T211433_
33 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:01
Finished backup at 10-SEP-22

Starting Control File and SPFILE Autobackup at 10-SEP-22
piece handle=/home/StudentFirst/DBST670/autobackup/2022_09_10/o1_mf_s_1115068474_kktfmv4s_.b
Finished Control File and SPFILE Autobackup at 10-SEP-22

RMAN> backup archivelog all;

Starting backup at 10-SEP-22
current log archived
using channel ORA_DISK_1
channel ORA_DISK_1: starting archived log backup set
channel ORA_DISK_1: specifying archived log(s) in backup set
input archived log thread=1 sequence=5 RECID=1 STAMP=1115068454
input archived log thread=1 sequence=6 RECID=2 STAMP=1115068473
input archived log thread=1 sequence=7 RECID=3 STAMP=1115068891
channel ORA_DISK_1: starting piece 1 at 10-SEP-22
channel ORA_DISK_1: finished piece 1 at 10-SEP-22
piece handle=/home/StudentFirst/DBST670/backupset/2022_09_10/o1_mf_annnn_TAG20220910T212131_
31 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:03
Finished backup at 10-SEP-22

Starting Control File and SPFILE Autobackup at 10-SEP-22
piece handle=/home/StudentFirst/DBST670/autobackup/2022_09_10/o1_mf_s_1115068894_kktg0yoq_.b
Finished Control File and SPFILE Autobackup at 10-SEP-22

RMAN>
```

